

```
In [1]: # ===== Package Check & Auto-Install =====
import importlib
import sys
import subprocess

def install_if_missing(pkg_name, import_name=None):
    import_name = import_name or pkg_name
    try:
        importlib.import_module(import_name)
        print(f"✓ {pkg_name} already installed")
    except ImportError:
        print(f"△ {pkg_name} not found - installing...")
        subprocess.check_call([sys.executable, "-m", "pip", "install", pkg_name])
        print(f"✓ Installed {pkg_name} successfully")

required_packages = [
    ("matplotlib", "matplotlib"),
    ("numpy", "numpy"),
    ("pandas", "pandas"),
    ("seaborn", "seaborn"),
    ("statsmodels", "statsmodels"),
    ("scikit-learn", "sklearn"),
    ("category_encoders", "category_encoders")
]

for pkg, import_name in required_packages:
    install_if_missing(pkg, import_name)

# ===== Imports =====

# hide warnings to keep printout clean
import warnings
warnings.filterwarnings("ignore")

import os
os.environ["PANDAS_IGNORE_BOTTLENECK"] = "1"
import numpy as np
import pandas as pd
import seaborn as sns
import math
```

```
import statsmodels.formula.api as smf
from category_encoders import TargetEncoder
import matplotlib.pyplot as plt

# sklearn
from sklearn.linear_model import LinearRegression, RidgeCV, LassoCV
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.model_selection import train_test_split, KFold, cross_val_score, GridSearchCV
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.tree import DecisionTreeRegressor
```

- ✓ matplotlib already installed
- ✓ numpy already installed
- ✓ pandas already installed
- ✓ seaborn already installed
- ✓ statsmodels already installed
- ✓ scikit-learn already installed
- ✓ category_encoders already installed

Data Cleaning

```
In [2]: df = pd.read_csv('data/rideshare_kaggle.csv')

# select variables needed for this project
selected_vars = [
    "id", "price", "distance", "cab_type", "name",
    "timestamp", "hour", "day", "source", "destination",
    "temperature", "precipIntensity", "precipProbability", "cloudCover",
    "surge_multiplier", "short_summary"
]

df = df[selected_vars].copy()
df.head()
```

Out[2]:

	id	price	distance	cab_type	name	timestamp	hour	day	source	destination	temperature	precipitation
0	424553bb-7174-41ea-aeb4-fe06d4f4b9d7	5.0	0.44	Lyft	Shared	1.544953e+09	9	16	Haymarket Square	North Station	42.34	
1	4bd23055-6827-41c6-b23b-3c491f24e74d	11.0	0.44	Lyft	Lux	1.543284e+09	2	27	Haymarket Square	North Station	43.58	
2	981a3613-77af-4620-a42a-0c0866077d1e	7.0	0.44	Lyft	Lyft	1.543367e+09	1	28	Haymarket Square	North Station	38.33	
3	c2d88af2-d278-4bfd-a8d0-29ca77cc5512	26.0	0.44	Lyft	Lux Black XL	1.543554e+09	4	30	Haymarket Square	North Station	34.38	
4	e0126e1f-8ca9-4f2e-82b3-50505a09db9a	9.0	0.44	Lyft	Lyft XL	1.543463e+09	3	29	Haymarket Square	North Station	37.44	

In [3]:

```
# rename the original day to day_of_month
df['day_of_month'] = df['day'].astype(int)

# fix timestamp and create proper day variables
df['datetime'] = pd.to_datetime(df['timestamp'], unit='s')

df["month_name"] = df["datetime"].dt.month_name() # January, February, ...
df["day_of_week"] = df["datetime"].dt.day_name() # Monday, Tuesday, ...

# create a new binary indicator for is_weekend
df['is_weekend'] = df['day_of_week'].isin(["Saturday", "Sunday"])

# remove variables not needed
df = df.drop(columns=["day", "timestamp", "datetime"])
```

`df.head()`

Out [3]:

	<code>id</code>	<code>price</code>	<code>distance</code>	<code>cab_type</code>	<code>name</code>	<code>hour</code>	<code>source</code>	<code>destination</code>	<code>temperature</code>	<code>precipIntensity</code>	<code>precipF</code>
0	424553bb-7174-41ea-aeb4-fe06d4f4b9d7	5.0	0.44	Lyft	Shared	9	Haymarket Square	North Station	42.34	0.0000	
1	4bd23055-6827-41c6-b23b-3c491f24e74d	11.0	0.44	Lyft	Lux	2	Haymarket Square	North Station	43.58	0.1299	
2	981a3613-77af-4620-a42a-0c0866077d1e	7.0	0.44	Lyft	Lyft	1	Haymarket Square	North Station	38.33	0.0000	
3	c2d88af2-d278-4bfd-a8d0-29ca77cc5512	26.0	0.44	Lyft	Lux Black XL	4	Haymarket Square	North Station	34.38	0.0000	
4	e0126e1f-8ca9-4f2e-82b3-50505a09db9a	9.0	0.44	Lyft	Lyft XL	3	Haymarket Square	North Station	37.44	0.0000	

In [4]:

`# General info
df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 693071 entries, 0 to 693070
Data columns (total 18 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   id               693071 non-null   object  
 1   price             637976 non-null   float64 
 2   distance          693071 non-null   float64 
 3   cab_type          693071 non-null   object  
 4   name              693071 non-null   object  
 5   hour              693071 non-null   int64  
 6   source             693071 non-null   object  
 7   destination        693071 non-null   object  
 8   temperature         693071 non-null   float64 
 9   precipIntensity    693071 non-null   float64 
 10  precipProbability  693071 non-null   float64 
 11  cloudCover         693071 non-null   float64 
 12  surge_multiplier   693071 non-null   float64 
 13  short_summary      693071 non-null   object  
 14  day_of_month       693071 non-null   int64  
 15  month_name         693071 non-null   object  
 16  day_of_week        693071 non-null   object  
 17  is_weekend         693071 non-null   bool    
dtypes: bool(1), float64(7), int64(2), object(8)
memory usage: 90.6+ MB
```

```
In [5]: # Shape
print(f"Rows: {df.shape[0]}, Columns: {df.shape[1]}")
```

```
Rows: 693071, Columns: 18
```

```
In [6]: # Check for duplicates
print(f"Duplicated rows: {df.duplicated().sum()}")
```

```
Duplicated rows: 0
```

```
In [7]: # number of missing values
df.isna().sum().sort_values(ascending=False).head()
```

```
Out[7]: price      55095
         id          0
        day_of_week    0
      month_name     0
   day_of_month     0
      dtype: int64
```

EDA

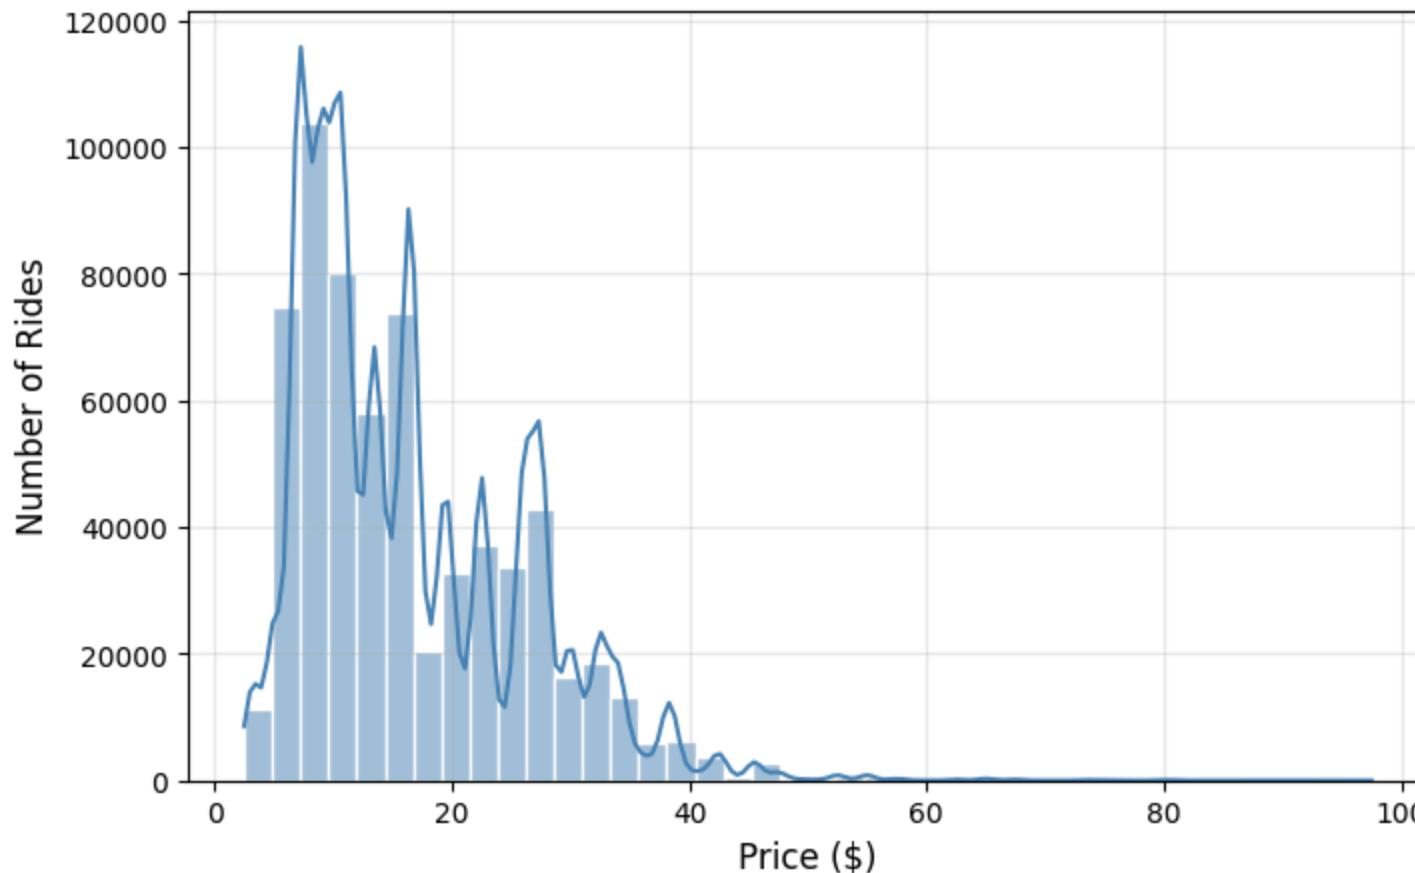
Outcome Variable Check

```
In [8]: # distribution of price (original), with na values removed
prices = df['price'].dropna()

plt.figure(figsize=(8,5))
sns.histplot(prices, bins=40, kde=True, color='steelblue', edgecolor='white')

plt.title('Distribution of Ride Prices', fontsize=14)
plt.xlabel('Price ($)', fontsize=12)
plt.ylabel('Number of Rides', fontsize=12)
plt.grid(alpha=0.3)
plt.show()
```

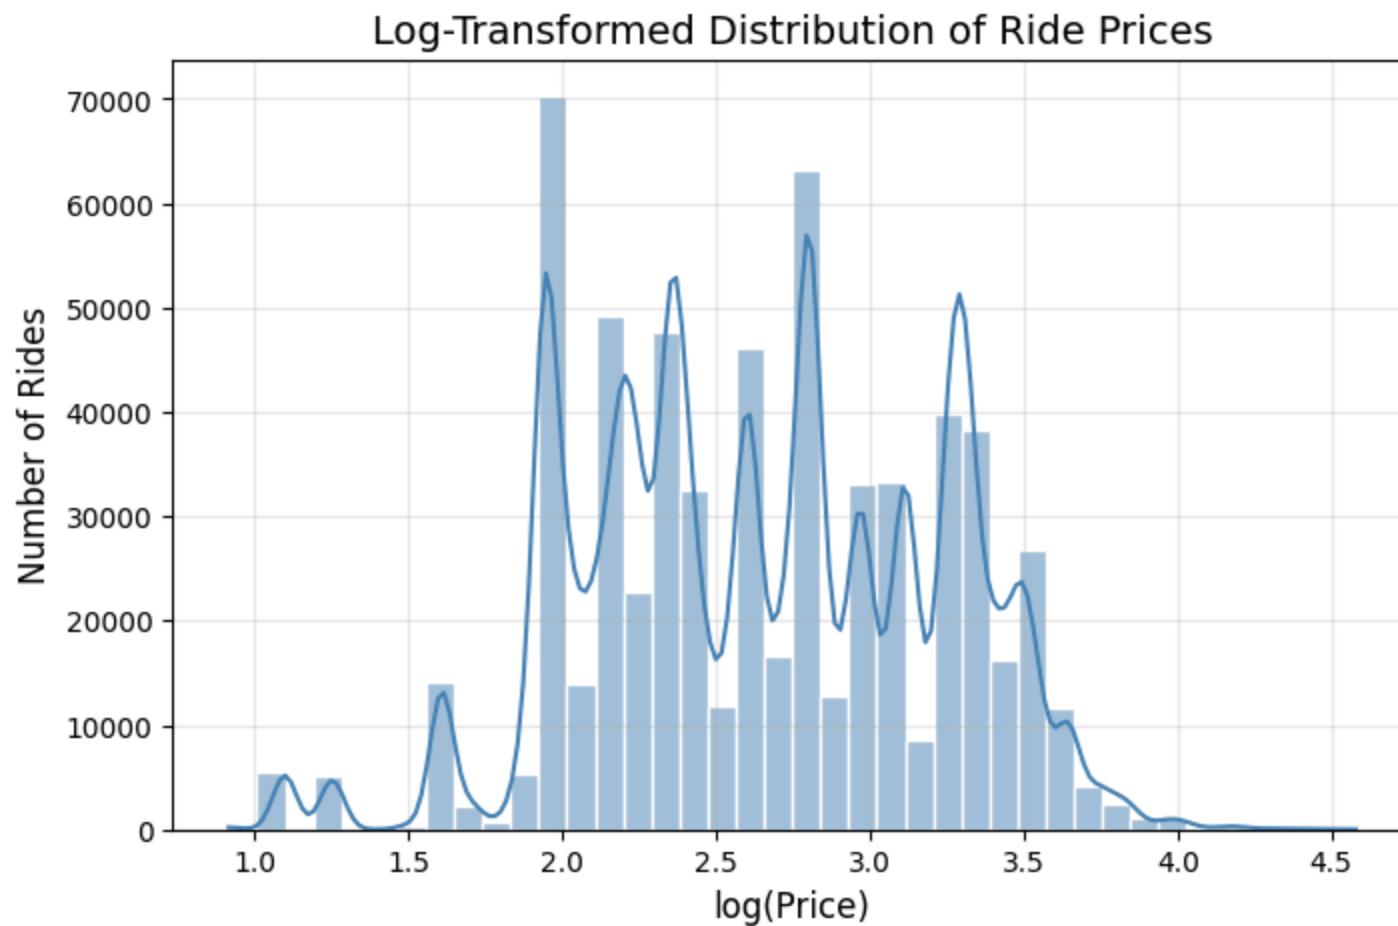
Distribution of Ride Prices



```
In [9]: # distribution of price (transformed), with na values removed
# Use np.log1p() to handle any zero or near-zero prices safely
prices = np.log(df['price'].dropna())

plt.figure(figsize=(8,5))
sns.histplot(prices, bins=40, kde=True, color='steelblue', edgecolor='white')

plt.title('Log-Transformed Distribution of Ride Prices', fontsize=14)
plt.xlabel('log(Price)', fontsize=12)
plt.ylabel('Number of Rides', fontsize=12)
plt.grid(alpha=0.3)
plt.show()
```



Categorical Variables Check

```
In [10]: categorical_var = [  
    "cab_type",  
    "name",  
    "source",  
    "destination",  
    "short_summary",  
    "month_name",  
    "day_of_week",
```

```

    "is_weekend"
]
```

```
In [11]: cat_summary = df[categorical_var].describe()

print("\n===== Categorical Variables Summary =====")
display(cat_summary)
```

===== Categorical Variables Summary =====

	cab_type	name	source	destination	short_summary	month_name	day_of_week	is_weekend
count	693071	693071	693071	693071	693071	693071	693071	693071
unique	2	13	12	12	9	2	7	2
top	Uber	UberXL	Financial District	Financial District	Overcast	December	Tuesday	False
freq	385663	55096	58857	58851	218895	406614	124949	511373

```
In [12]: n = len(categorical_var)
rows = math.ceil(n / 3)

fig, axes = plt.subplots(rows, 3, figsize=(22, 6 * rows))
axes = axes.flatten()

for i, col in enumerate(categorical_var):
    # group category by price missing / not missing
    counts = df.groupby([col, df['price'].isna()]).size().unstack(fill_value=0)
    # rename columns for clarity
    counts.columns = ['price_not_missing', 'price_missing']

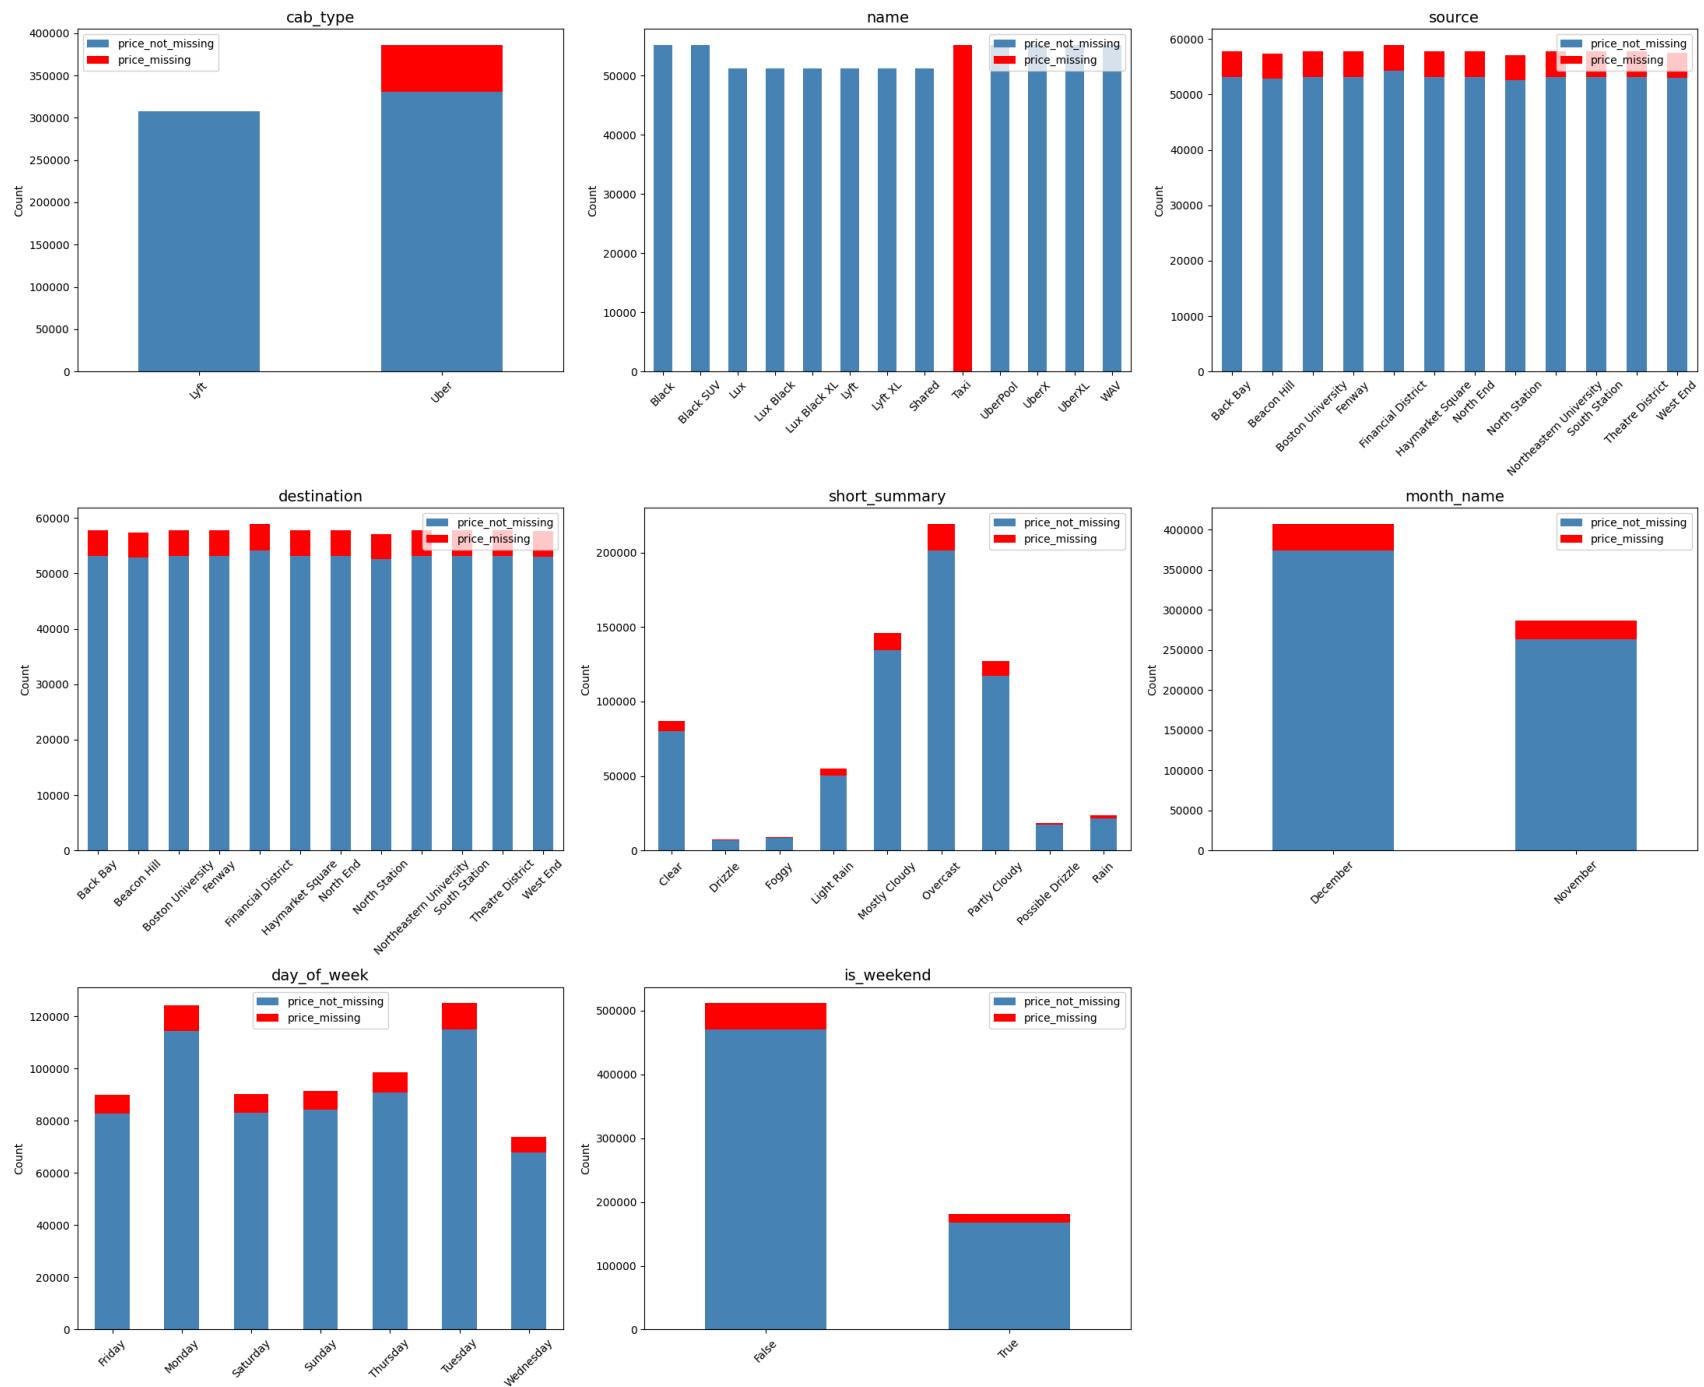
    # stacked bar plot
    counts[['price_not_missing', 'price_missing']].plot(
        kind='bar',
        stacked=True,
        ax=axes[i],
        color=['steelblue', 'red'] # normal + missing price in red
    )

    axes[i].set_title(f"{col}", fontsize=14)
    axes[i].set_xlabel("")
    axes[i].set_ylabel("Count")
```

```
axes[i].tick_params(axis='x', labelrotation=45)

# hide unused subplots if number of vars % 3 != 0
for j in range(i + 1, len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()
```



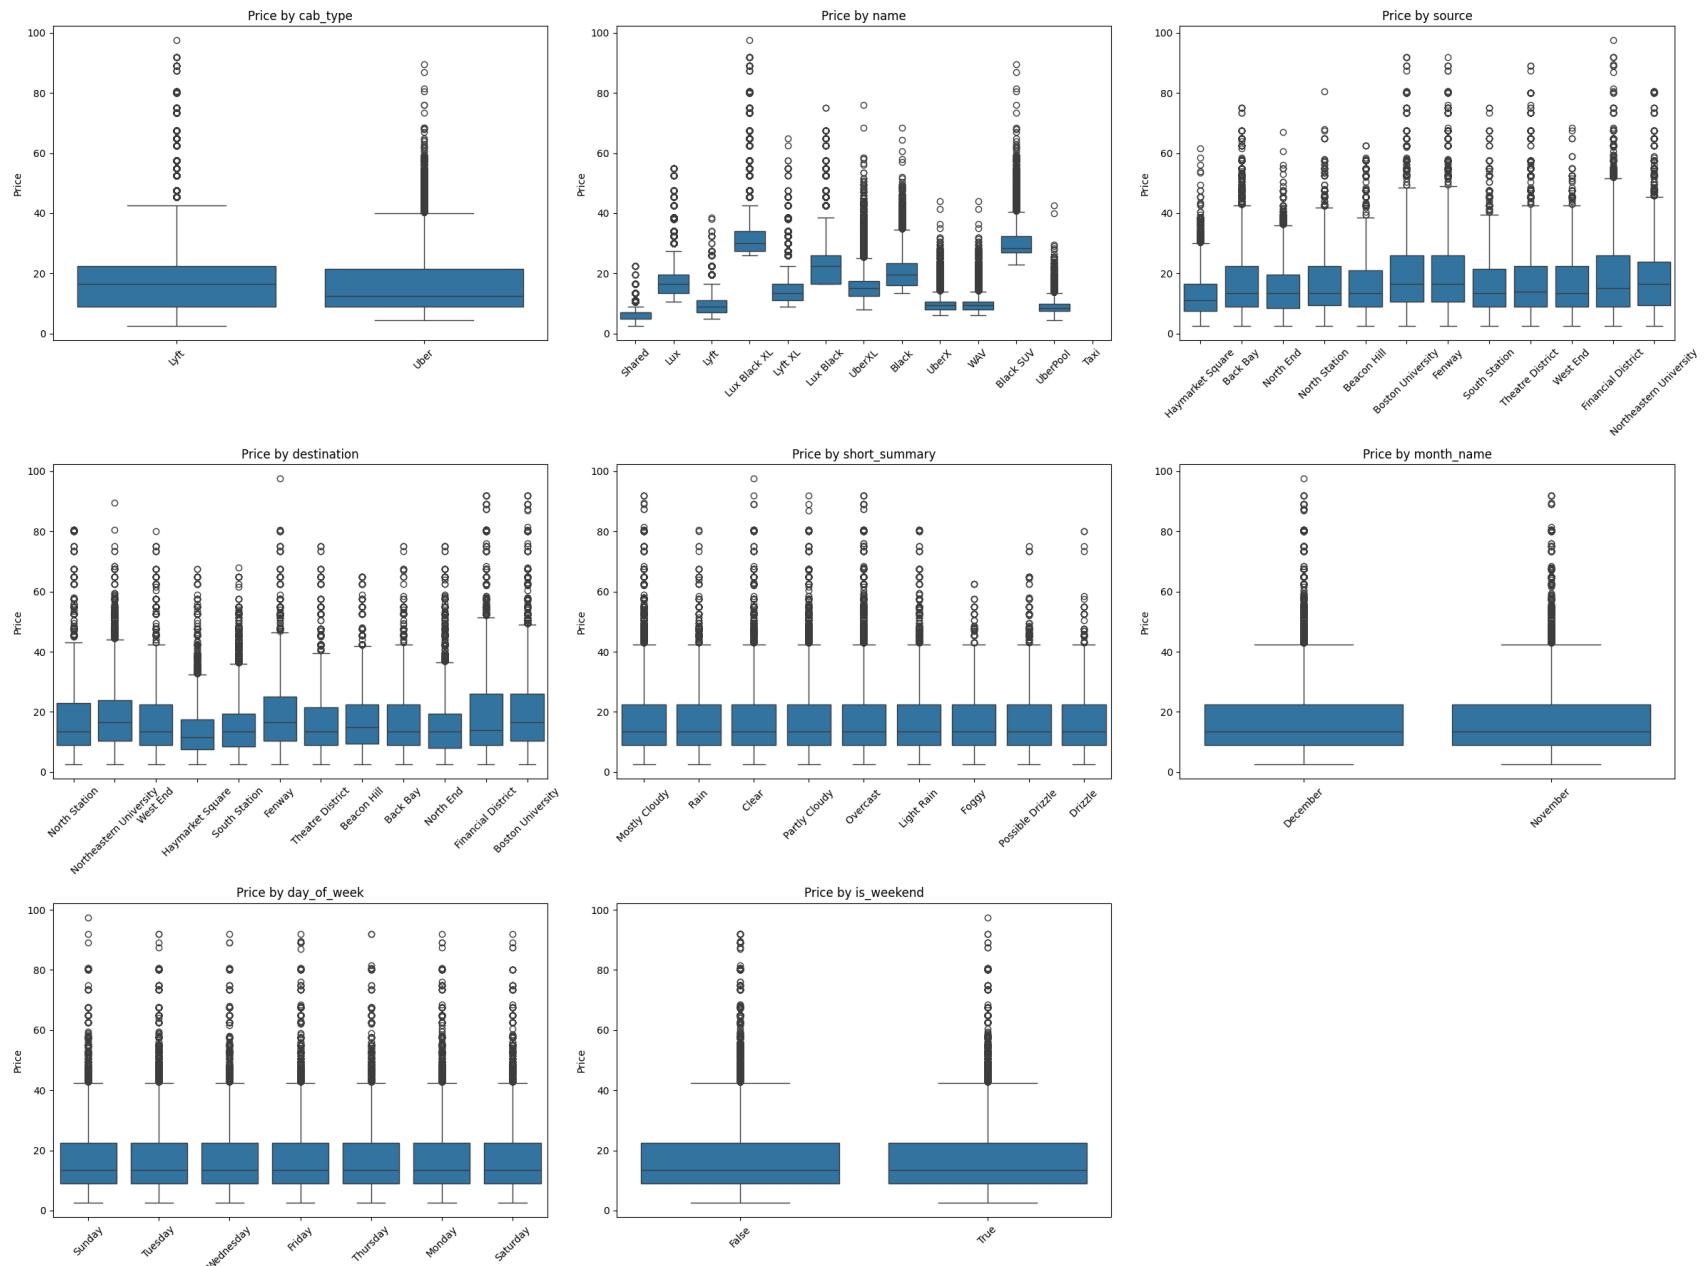
```
In [13]: n = len(categorical_var)
rows = math.ceil(n / 3)

fig, axes = plt.subplots(rows, 3, figsize=(24, 6 * rows))
axes = axes.flatten()

for i, col in enumerate(categorical_var):
    sns.boxplot(data=df, x=col, y="price", ax=axes[i])
    axes[i].set_title(f"Price by {col}")
    axes[i].set_xlabel("")
    axes[i].set_ylabel("Price")
    axes[i].tick_params(axis="x", rotation=45)

# hide unused subplots
for j in range(i + 1, len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()
```



Across categorical predictors, price exhibits relatively consistent patterns with a few meaningful differences. Cab type and service name show the strongest price variation, with premium services (e.g., Lyft Lux, Lyft XL) generally yielding higher fares compared to standard offerings. Pickup and drop-off locations also influence pricing, with certain hubs (e.g., airports and

major transport stations) associated with higher median fares, likely reflecting longer distances and greater demand. In contrast, weather-related categories (short_summary), day of week, weekend indicator, and month show only modest differences in price distributions, suggesting they play a secondary role relative to service level and location. Overall, the categorical EDA indicates that ride type and geographic endpoints are key drivers of price, while temporal and weather-based variables may have weaker direct effects.

Numerical Variables Check

```
In [14]: numerical_var = [  
    "distance",  
    "hour",  
    "temperature",  
    "precipIntensity",  
    "precipProbability",  
    "cloudCover",  
    "surge_multiplier",  
    "day_of_month"  
]
```

```
In [15]: num_summary = df[numerical_var].describe().round(2)  
  
print("\n===== Numerical Variables Summary =====")  
display(num_summary)
```

===== Numerical Variables Summary =====

	distance	hour	temperature	precipIntensity	precipProbability	cloudCover	surge_multiplier	day_of_month
count	693071.00	693071.00	693071.00	693071.00	693071.00	693071.00	693071.00	693071.00
mean	2.19	11.62	39.58	0.01	0.15	0.69	1.01	17.79
std	1.14	6.95	6.73	0.03	0.33	0.36	0.09	9.98
min	0.02	0.00	18.91	0.00	0.00	0.00	1.00	1.00
25%	1.28	6.00	36.45	0.00	0.00	0.37	1.00	13.00
50%	2.16	12.00	40.49	0.00	0.00	0.82	1.00	17.00
75%	2.92	18.00	43.58	0.00	0.00	1.00	1.00	28.00
max	7.86	23.00	57.22	0.14	1.00	1.00	3.00	30.00

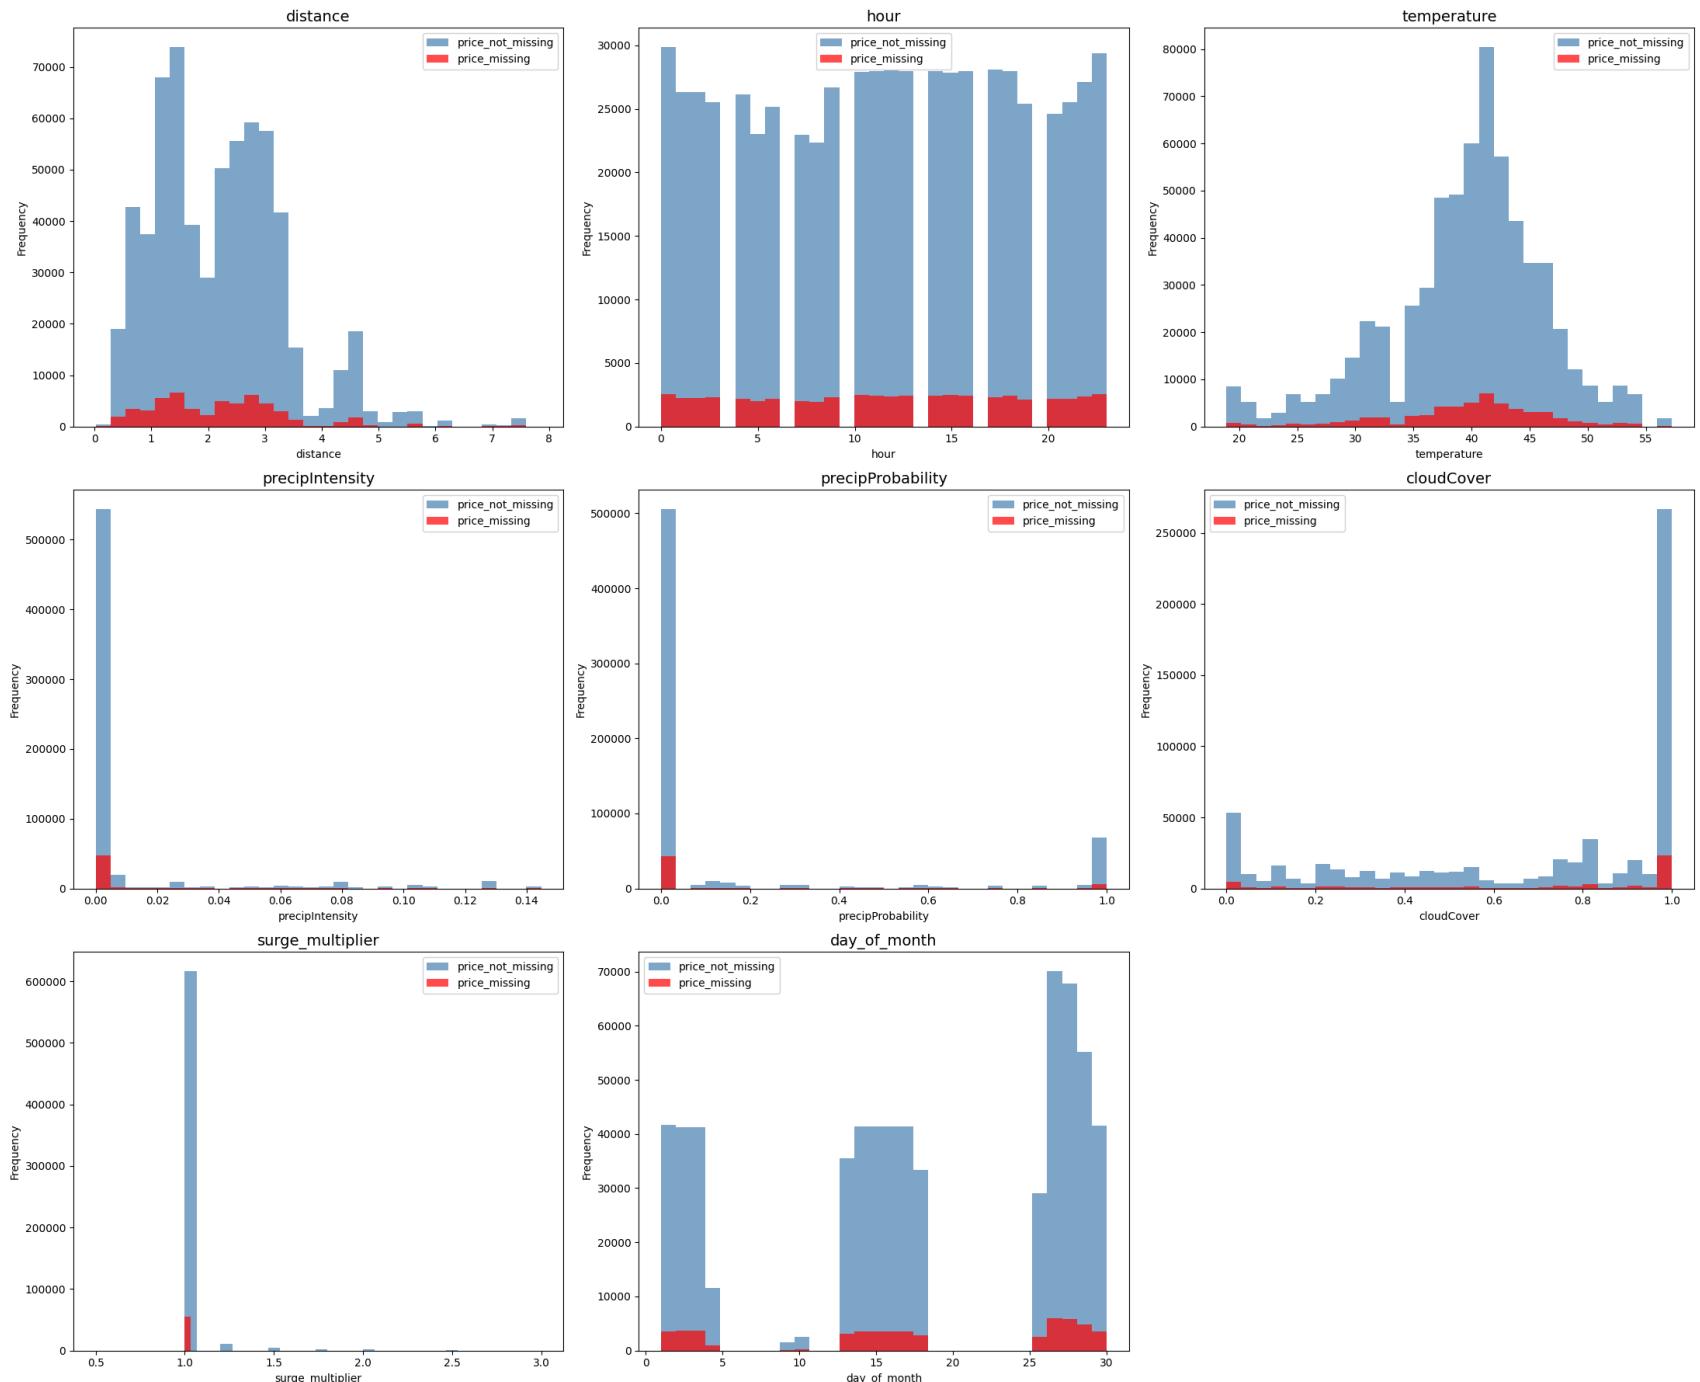
```
In [16]: n = len(numerical_var)
rows = math.ceil(n / 3)

fig, axes = plt.subplots(rows, 3, figsize=(22, 6 * rows))
axes = axes.flatten()

for i, col in enumerate(numerical_var):
    # valid price values
    df[df["price"].notna()][col].plot(
        kind='hist',
        bins=30,
        alpha=0.7,
        color='steelblue',
        ax=axes[i],
        label="price_not_missing"
    )

    # missing price values
    df[df["price"].isna()][col].plot(
        kind='hist',
        bins=30,
        alpha=0.7,
        color='red',
        ax=axes[i],
        label="price_missing"
```

```
)  
  
axes[i].set_title(f"{col}", fontsize=14)  
axes[i].set_xlabel(col)  
axes[i].set_ylabel("Frequency")  
axes[i].legend()  
  
# hide unused subplots (if number of vars not divisible by 3)  
for j in range(i + 1, len(axes)):  
    fig.delaxes(axes[j])  
  
plt.tight_layout()  
plt.show()
```



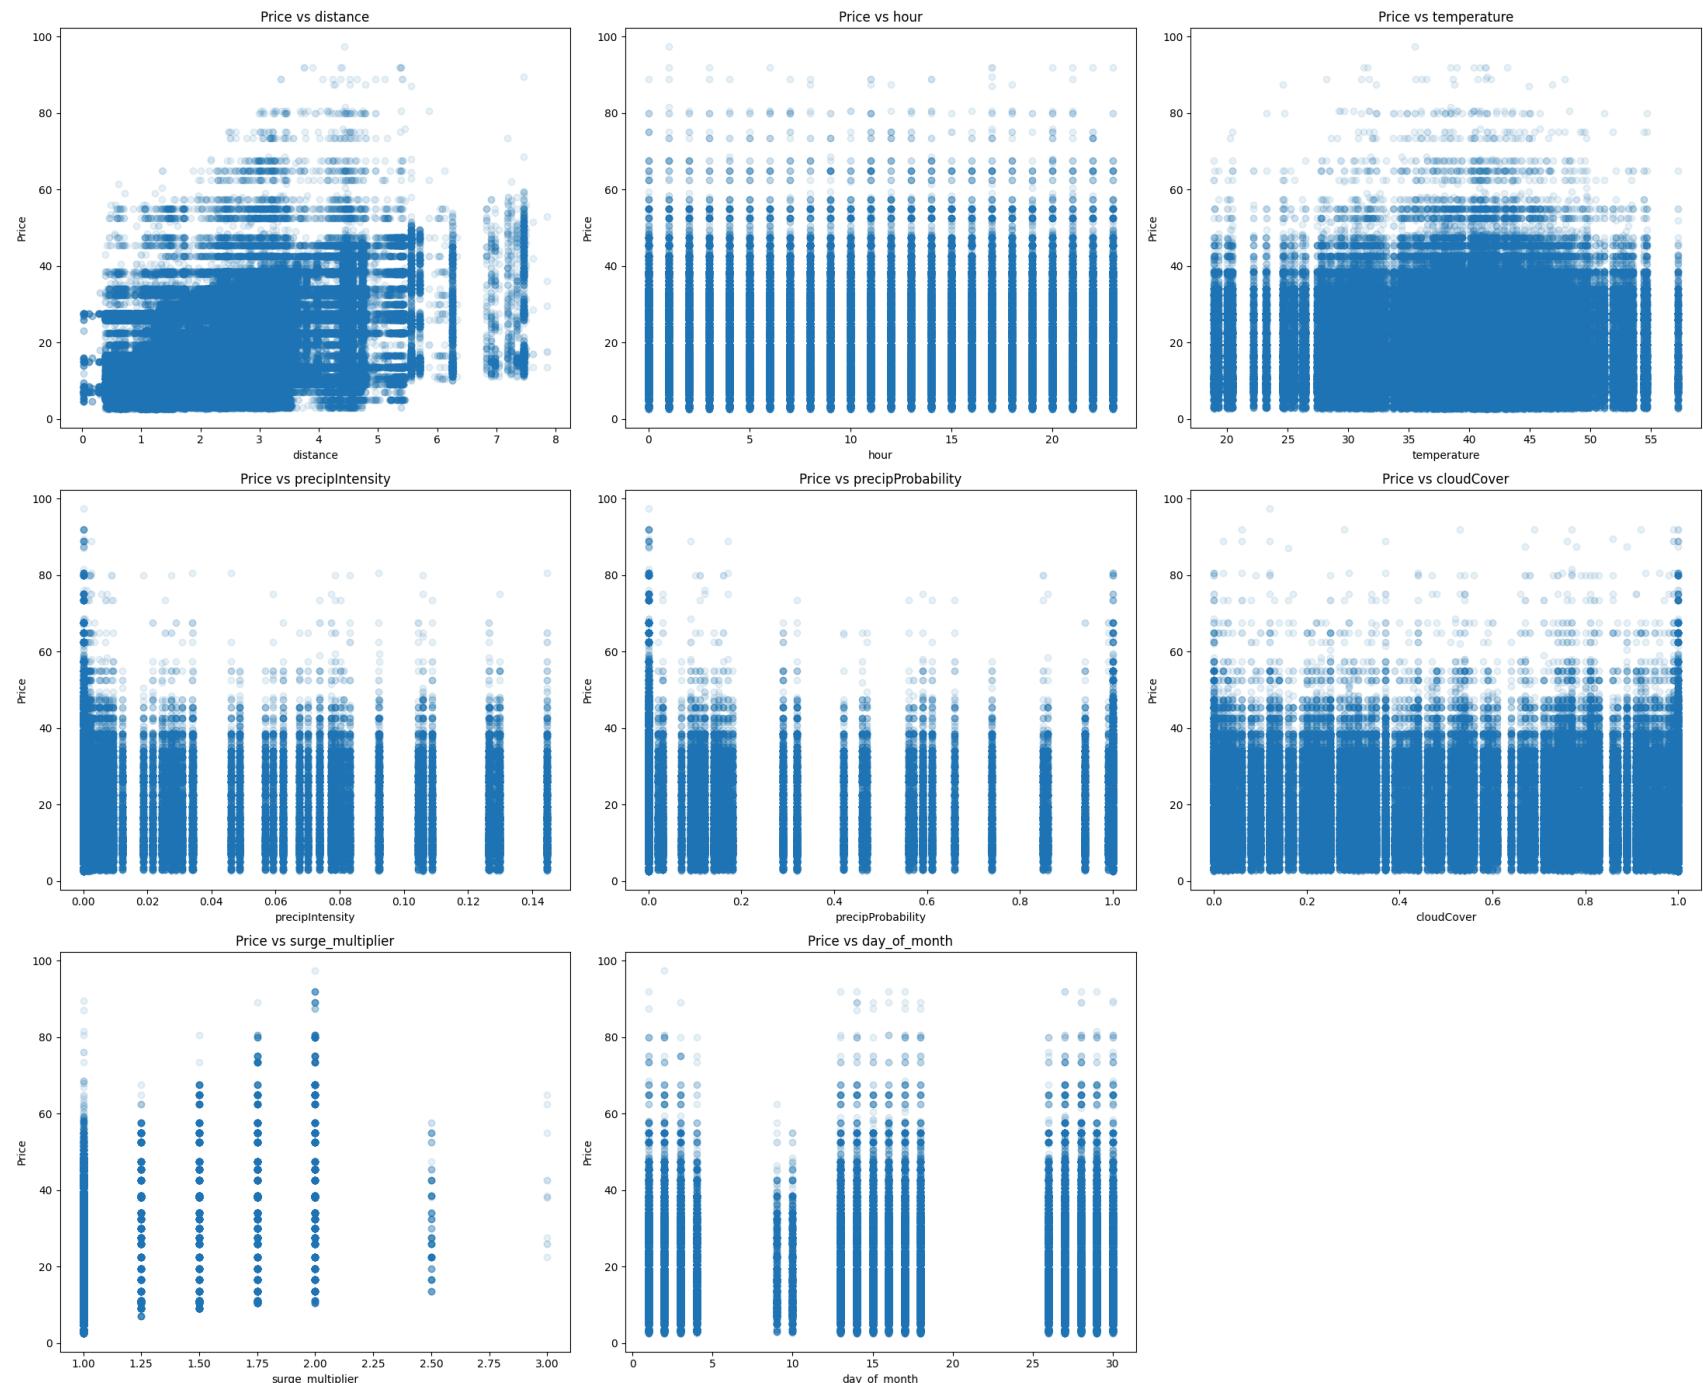
```
In [17]: n = len(numerical_var)
rows = math.ceil(n / 3)

fig, axes = plt.subplots(rows, 3, figsize=(22, 6 * rows))
axes = axes.flatten()

for i, col in enumerate(numerical_var):
    axes[i].scatter(df[col], df["price"], alpha=0.1)
    axes[i].set_title(f"Price vs {col}")
    axes[i].set_xlabel(col)
    axes[i].set_ylabel("Price")

# hide unused subplots if not multiple of 3
for j in range(i + 1, len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()
```



```
In [18]: # Outlier detection
outlier_summary = {}

for col in numerical_var:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1
    lower = Q1 - 1.5 * IQR
    upper = Q3 + 1.5 * IQR

    outliers = ((df[col] < lower) | (df[col] > upper)).sum()
    outlier_summary[col] = outliers

pd.DataFrame.from_dict(outlier_summary, orient='index', columns=['Outlier Count'])
```

Out[18]:

Outlier Count

distance	8662
hour	0
temperature	36659
precipIntensity	150828
precipProbability	150828
cloudCover	0
surge_multiplier	20975
day_of_month	0

```
In [19]: # Scaling
scaler = StandardScaler()
scaled = scaler.fit_transform(df[numerical_var])

scaled_df = pd.DataFrame(scaled, columns=[c + '_scaled' for c in numerical_var])
scaled_df.head()
```

Out[19]:

	distance_scaled	hour_scaled	temperature_scaled	precipIntensity_scaled	precipProbability_scaled	cloudCover_scaled
0	-1.536021	-0.376957	0.409691	-0.331672	-0.444083	0.093430
1	-1.536021	-1.384425	0.594048	4.497229	2.596003	0.874388
2	-1.536021	-1.528349	-0.186496	-0.331672	-0.444083	-1.831073
3	-1.536021	-1.096577	-0.773763	-0.331672	-0.444083	-1.914747
4	-1.536021	-1.240501	-0.318817	-0.331672	-0.444083	-0.687527

Findings: The numerical predictors in our dataset exhibit a wide range of distributional behaviors.

- Distance shows a multimodal distribution reflecting different ride types (short local trips vs. longer airport rides).
- Hour displays usage across all 24 hours, with moderately higher demand during daytime and early evening.
- CloudCover shows a bimodal distribution, heavily concentrated at both 0 (clear) and 1 (fully overcast), with fewer values in between.
- Temperature is moderately distributed.
- Weather variables, especially precipIntensity and precipProbability, are extremely right-skewed, with over 150,000 outliers each, reflecting the rarity of heavy precipitation events. CloudCover is left-skewed.
- Surge_multiplier is highly concentrated at 1.0, indicating that the vast majority of rides occur without surge pricing.
- Day_of_month has three peaks, suggesting limited temporal coverage in the dataset.

Scatterplots further confirm that distance is the only predictor showing a visible positive relationship with price. Other variables exhibit weak or no linear associations.

All predictors were standardized.

Correlation check

In [20]:

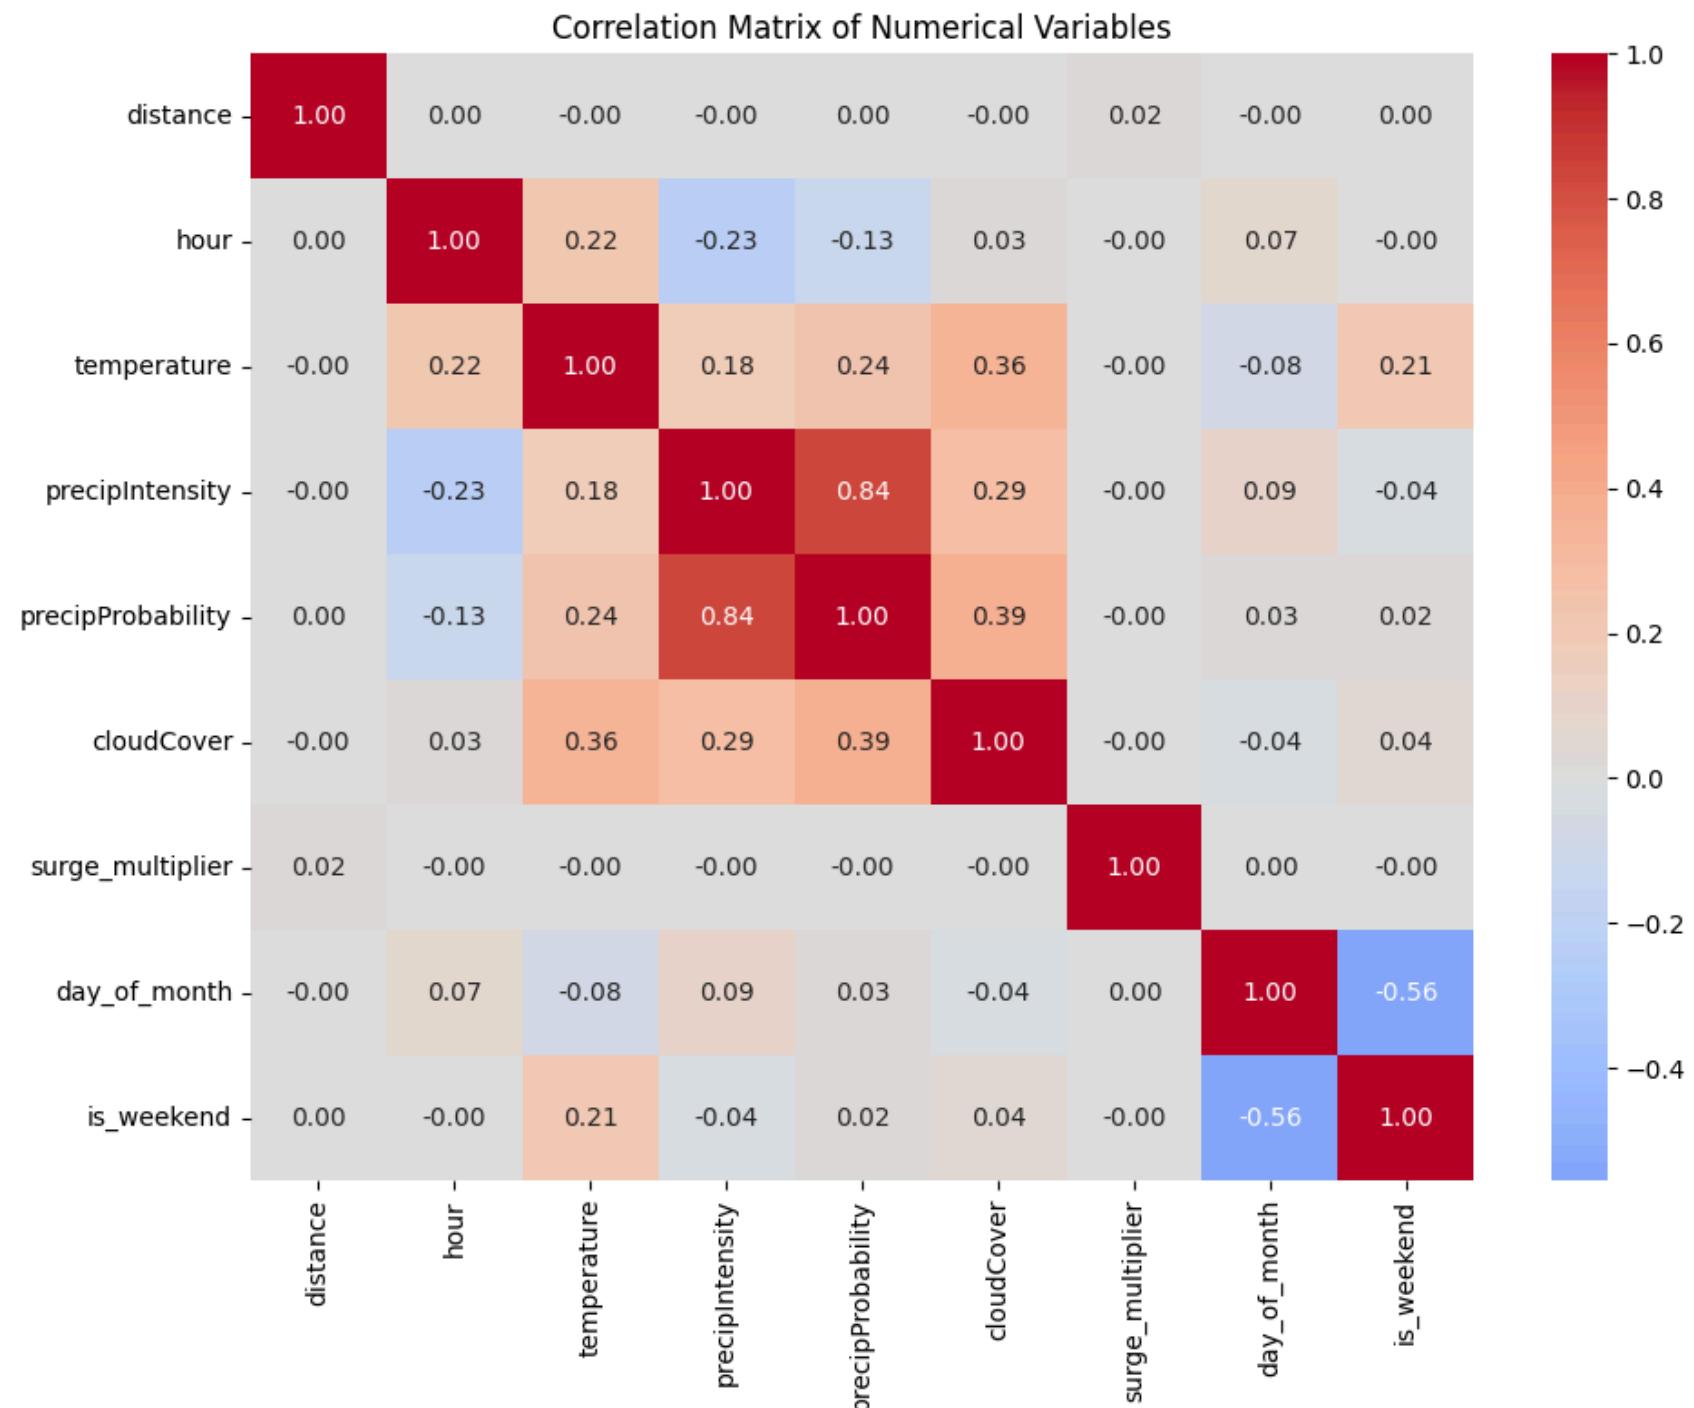
```
corr_var = [
    "distance",
    "hour",
    "temperature",
    "precipIntensity",
    "precipProbability",
```

```
"cloudCover",
"surge_multiplier",
"day_of_month",
"is_weekend"
]

corr_matrix = df[corr_var].corr()

plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, fmt='.2f', cmap='coolwarm', center=0)
plt.title('Correlation Matrix of Numerical Variables')
plt.tight_layout()
plt.show()

print("Correlation Matrix:")
print(corr_matrix)
```



Correlation Matrix:

	distance	hour	temperature	precipIntensity	\
distance	1.000000	0.002280	-0.002884	-0.000256	
hour		1.000000	0.218769	-0.233349	
temperature			1.000000	0.182724	
precipIntensity				1.000000	
precipProbability					0.838470
cloudCover					0.288960
surge_multiplier					-0.001530
day_of_month					0.091664
is_weekend					-0.037736
	precipProbability	cloudCover	surge_multiplier	\	
distance	0.000371	-0.000905	0.024769		
hour		-0.129725	0.026525	-0.000077	
temperature			0.355156	-0.001572	
precipIntensity				-0.001530	
precipProbability					0.387114
cloudCover					-0.002103
surge_multiplier					1.000000
day_of_month					0.001389
is_weekend					-0.001320
	day_of_month	is_weekend			
distance	-0.000627	0.001523			
hour		0.066090	-0.002531		
temperature			0.212279		
precipIntensity				-0.037736	
precipProbability					0.030903
cloudCover					-0.041190
surge_multiplier					0.001389
day_of_month					1.000000
is_weekend					-0.555498

Insights:

- `precipIntensity` and `precipProbability` are highly correlated ($r = 0.84$).

This is reasonable because higher rainfall intensity typically occurs when the probability of rain is also high—both variables capture related aspects of precipitation events.

We will drop `precipProbability` and use `precipIntensity` for modeling.

- `temperature` and `cloudCover` show moderate correlation ($r = 0.36$).

This is reasonable because cloud cover influences temperature patterns—clouds trap heat at night and reduce warming during the day—creating a natural meteorological relationship.

We will drop `cloudCover` and use `temperature` for modeling.

- `name` and `cab_type` are perfectly correlated ($r = 1$).

This is expected because `name` (the specific ride type) is nested within `cab_type` (Uber or Lyft), so both variables encode the same information.

We will drop `name` and use `cab_type` for modeling.

- `is_weekend` and `day_of_month` are negatively correlated ($r = -0.56$). This is expected because `is_weekend` is a binary indicator created from `day_of_month`. We will drop `day_of_month` and use `is_weekend` for modeling.

- Most remaining variables show weak correlations ($|r| < 0.4$), indicating minimal multicollinearity concerns among the retained predictors.

final features selection: `feature_cols = ['distance', 'surge_multiplier', 'temperature', 'precipIntensity', 'hour', 'is_weekend', 'month_name', 'source', 'destination', 'name', 'cab_type']`

Baseline model: $\log(\text{price}) \sim \text{distance}$

```
In [21]: # train-test-validation split

# remove missing values
df_clean = df.dropna()

X = df_clean[['distance']].values
y = np.log(df_clean['price']) # log-transform price

# first split: 80% train+val, 20% test
X_train_and_val, X_test, y_train_and_val, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# second split: 75% of train+val for train (60% overall), 25% for val (20% overall)
X_train, X_val, y_train, y_val = train_test_split(
```

```
X_train_and_val, y_train_and_val, test_size=0.25, random_state=42
)

print(f"\nData split sizes:")
print(f"Train: {len(X_train)} ({len(X_train)/len(df_clean)*100:.1f}%)")
print(f"Validation: {len(X_val)} ({len(X_val)/len(df_clean)*100:.1f}%)")
print(f"Test: {len(X_test)} ({len(X_test)/len(df_clean)*100:.1f}%)")
print(f"Train+Val: {len(X_train_and_val)} ({len(X_train_and_val)/len(df_clean)*100:.1f}%)")
```

Data split sizes:
Train: 382785 (60.0%)
Validation: 127595 (20.0%)
Test: 127596 (20.0%)
Train+Val: 510380 (80.0%)

```
In [22]: baseline_model = LinearRegression()
baseline_model.fit(X_train, y_train)

y_train_pred = baseline_model.predict(X_train)
y_val_pred = baseline_model.predict(X_val)

train_r2 = r2_score(y_train, y_train_pred)
val_r2 = r2_score(y_val, y_val_pred)
train_rmse = np.sqrt(mean_squared_error(y_train, y_train_pred))
val_rmse = np.sqrt(mean_squared_error(y_val, y_val_pred))

print("\nBaseline Model: log(price) ~ distance")
print(f"Coefficient (distance): {baseline_model.coef_[0]:.4f}")
print(f"Intercept: {baseline_model.intercept_:.4f}")
print(f"\nTrain R-sq: {train_r2:.4f}")
print(f"Validation R-sq: {val_r2:.4f}")
print(f"Train RMSE: {train_rmse:.4f}")
print(f"Validation RMSE: {val_rmse:.4f}")

# Visualize
plt.figure(figsize=(10, 6))
plt.scatter(X_train, y_train, alpha=0.3, label='Training data', s=1)

X_train_sorted = np.sort(X_train, axis=0)
plt.plot(X_train_sorted,
         baseline_model.predict(X_train_sorted),
         color='red', linewidth=2, label='Fitted line')
```

```
plt.xlabel('Distance (miles)')
plt.ylabel('log(Price)')
plt.title('Baseline Model: log(Price) vs Distance')
plt.legend()
plt.tight_layout()
plt.show()
```

Baseline Model: $\log(\text{price}) \sim \text{distance}$

Coefficient (distance): 0.1685

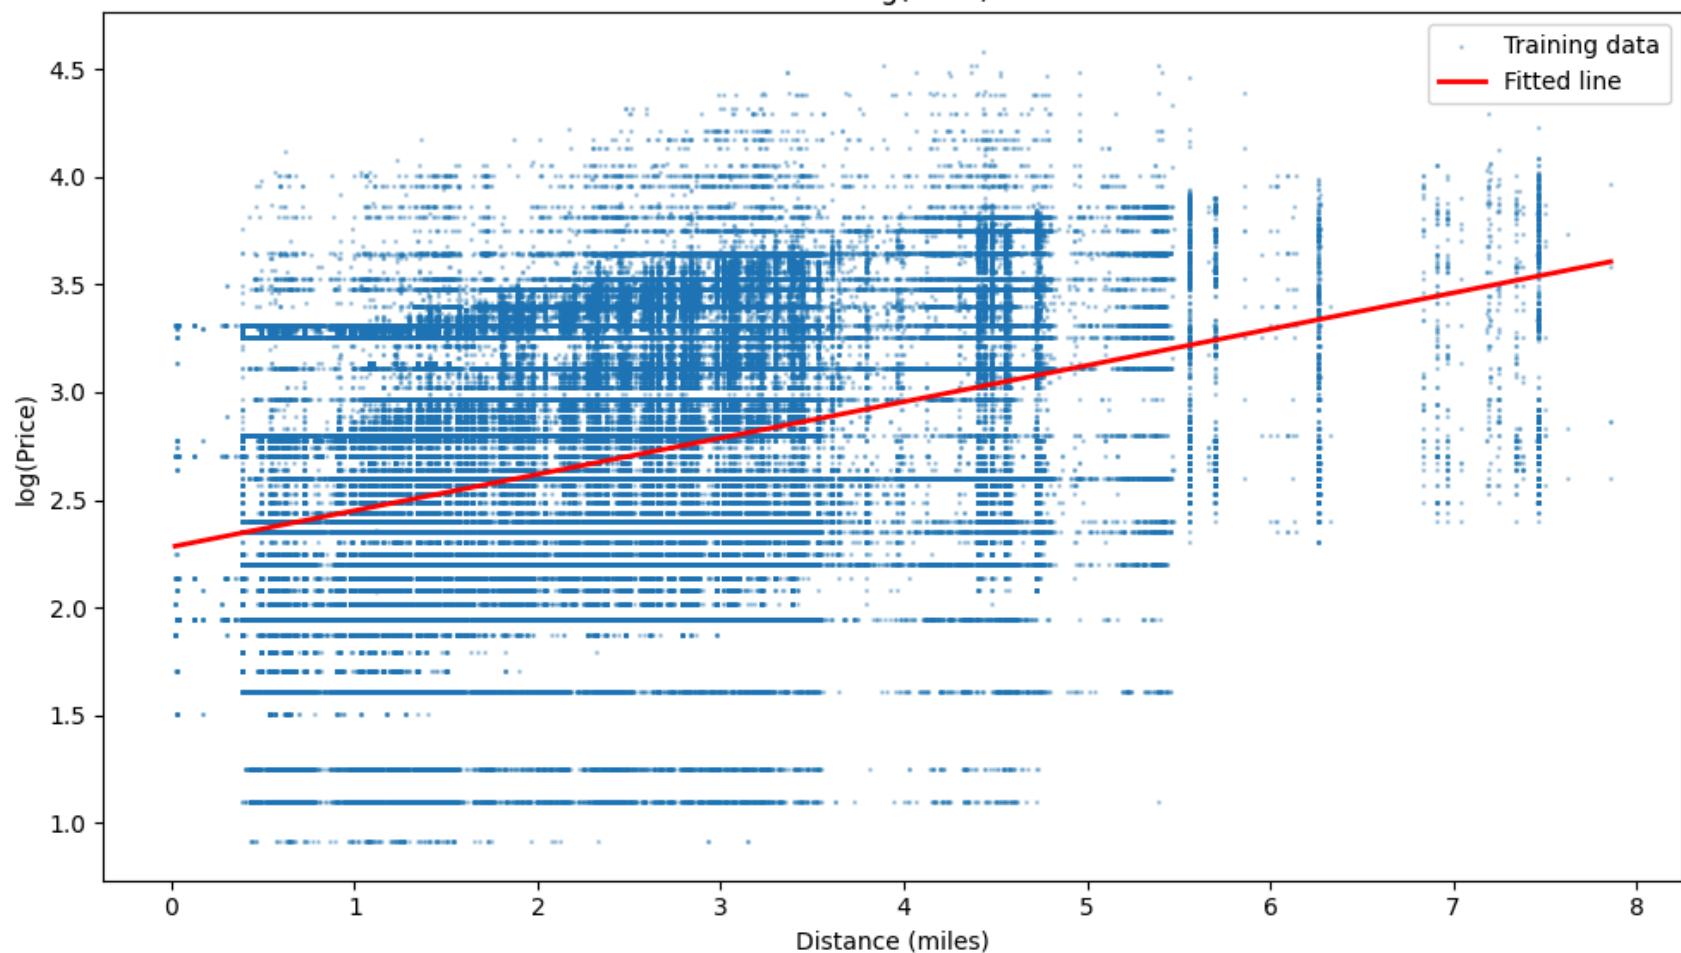
Intercept: 2.2819

Train R-sq: 0.1135

Validation R-sq: 0.1169

Train RMSE: 0.5347

Validation RMSE: 0.5368

Baseline Model: $\log(\text{Price})$ vs Distance

The baseline model shows a positive linear relationship between distance and log-transformed price (coefficient = 0.169), but distance alone explains only about 11-12% of the variance in pricing (Train R^2 = 0.114, Validation R^2 = 0.117). The distinct horizontal banding in the scatter plot and low R^2 suggest that other factors—such as service type, surge pricing, and contextual variables—are necessary to adequately capture the ride-sharing pricing structure.

Multiple Linear Regression

Target encoding

We employ **target encoding** for the `source` and `destination` neighborhood variables rather than traditional one-hot encoding.

Our dataset contains 12 unique pickup neighborhoods (`source`) and 12 unique dropoff neighborhoods (`destination`). One-hot encoding would create 24 binary indicator variables. Target encoding reduces these 24 dimensions to just 2 continuous variables. It replaces each neighborhood with its average $\log(\text{price})$ in the training set. This approach can efficiently encodes the economic characteristics of neighborhoods such as demand patterns, typical trip lengths, and surge likelihood into a single numerical value.

While we apply target encoding to locations, we deliberately use one-hot encoding for `name` (service class) and `cab_type` (platform). This hybrid approach allows us to create interpretable interaction terms such as `distance × UberXL` to capture how premium services charge different per-mile rates.

Target encoding represents an advanced feature engineering technique beyond the standard methods covered in our coursework. It demonstrates practical solutions to high-cardinality categorical variables, a common challenge in real-world machine learning applications.

```
In [23]: df_clean = df.dropna()

feature_cols = [
    'distance', 'surge_multiplier', 'temperature', 'precipIntensity',
    'hour', 'is_weekend', 'month_name',
    'source', 'destination', 'name', 'cab_type'
]

# feature matrix and target variable (used for both MLR and Decision Tree)
X_features = df_clean[feature_cols].copy()
y_target = np.log(df_clean['price']).copy()

X_train_val, X_test, y_train_val, y_test = train_test_split(
    X_features, y_target, test_size=0.2, random_state=42
)
```

```
X_train, X_val, y_train, y_val = train_test_split(
    X_train_val, y_train_val, test_size=0.25, random_state=42
)

# reset indices
X_train = X_train.reset_index(drop=True)
X_val = X_val.reset_index(drop=True)
X_test = X_test.reset_index(drop=True)
y_train = y_train.reset_index(drop=True)
y_val = y_val.reset_index(drop=True)
y_test = y_test.reset_index(drop=True)

print(f" Training set: {X_train.shape[0]}:,} samples ({X_train.shape[0]/len(df_clean)*100:.1f}%)")
print(f" Validation set: {X_val.shape[0]}:,} samples ({X_val.shape[0]/len(df_clean)*100:.1f}%)")
print(f" Test set: {X_test.shape[0]}:,} samples ({X_test.shape[0]/len(df_clean)*100:.1f}%)")
```

Training set: 382,785 samples (60.0%)
Validation set: 127,595 samples (20.0%)
Test set: 127,596 samples (20.0%)

In [24]:

```
# define columns for target encoding
location_cols = ['source', 'destination']
target_encoder = TargetEncoder(cols=location_cols, smoothing=1.0)

# fit on training data only
target_encoder.fit(X_train[location_cols], y_train)

# transform all sets
X_train_loc_encoded = target_encoder.transform(X_train[location_cols])
X_val_loc_encoded = target_encoder.transform(X_val[location_cols])
X_test_loc_encoded = target_encoder.transform(X_test[location_cols])

# rename columns
X_train_loc_encoded.columns = ['source_encoded', 'destination_encoded']
X_val_loc_encoded.columns = ['source_encoded', 'destination_encoded']
X_test_loc_encoded.columns = ['source_encoded', 'destination_encoded']
```

In [25]:

```
# mapping
# source neighborhoods
encoding_map_source = pd.DataFrame({
    'source': X_train['source'].values,
    'source_encoded': X_train_loc_encoded['source_encoded'].values
```

```
}).drop_duplicates().sort_values('source_encoded', ascending=False)

print("\nTop 5 most expensive pickup neighborhoods:")
print(encoding_map_source.head())
print("\nTop 5 least expensive pickup neighborhoods:")
print(encoding_map_source.tail())

# destination neighborhoods
encoding_map_dest = pd.DataFrame({
    'destination': X_train['destination'].values,
    'destination_encoded': X_train_loc_encoded['destination_encoded'].values
}).drop_duplicates().sort_values('destination_encoded', ascending=False)

print("\nTop 5 most expensive dropoff neighborhoods:")
print(encoding_map_dest.head())
print("\nTop 5 least expensive dropoff neighborhoods:")
print(encoding_map_dest.tail())
```

Top 5 most expensive pickup neighborhoods:

	source	source_encoded
2	Boston University	2.781737
19	Fenway	2.762028
9	Northeastern University	2.734369
7	Financial District	2.703400
1	Theatre District	2.659597

Top 5 least expensive pickup neighborhoods:

	source	source_encoded
10	Back Bay	2.624302
13	Beacon Hill	2.622324
4	South Station	2.607104
33	North End	2.575637
25	Haymarket Square	2.457238

Top 5 most expensive dropoff neighborhoods:

	destination	destination_encoded
8	Boston University	2.793605
0	Fenway	2.749527
1	Northeastern University	2.737349
4	Financial District	2.701635
25	North Station	2.657249

Top 5 least expensive dropoff neighborhoods:

	destination	destination_encoded
5	West End	2.634631
2	Theatre District	2.631630
13	North End	2.554999
11	South Station	2.547252
3	Haymarket Square	2.507349

One-hot encoding

```
In [26]: categorical_cols = ['name', 'cab_type', 'month_name']
# create dummy variables
X_train_categorical = pd.get_dummies(X_train[categorical_cols], drop_first=True, dtype=int)
X_val_categorical = pd.get_dummies(X_val[categorical_cols], drop_first=True, dtype=int)
X_test_categorical = pd.get_dummies(X_test[categorical_cols], drop_first=True, dtype=int)
```

```
all_categorical_cols = X_train_categorical.columns

# add missing columns
for col in all_categorical_cols:
    if col not in X_val_categorical.columns:
        X_val_categorical[col] = 0
    if col not in X_test_categorical.columns:
        X_test_categorical[col] = 0

# reorder columns
X_val_categorical = X_val_categorical[all_categorical_cols]
X_test_categorical = X_test_categorical[all_categorical_cols]
```

Combine all features

Our multiple linear regression model includes features organized into categories that capture different aspects of ride-sharing pricing, selected based on theoretical expectations and insights from our EDA.

Core pricing drivers

`Distance` serves as our primary explanatory variable since trip length theoretically determines ride-sharing costs. We include `surge_multiplier` to capture dynamic pricing adjustments during high demand periods. The service class variable (`name`) is essential because our EDA revealed that premium services like Lyft Lux have substantially higher prices than standard options, indicating different base rates and per-mile charges. The `cab_type` variable (Uber vs. Lyft) allows us to test whether the two platforms have systematically different pricing structures.

Spatial and temporal context

Pickup and dropoff locations (`source` and `destination`) influence pricing through neighborhood-specific demand patterns and typical trip characteristics. We apply target encoding to efficiently capture these location effects while avoiding dimensionality issues from one-hot encoding 24 neighborhoods. For temporal patterns, we include `hour` to capture rush hour effects, `is_weekend` to distinguish weekday commuting from weekend leisure travel, and `month` to account for seasonal variation in demand and pricing.

Weather

We include `temperature` and `precipIntensity` to capture how adverse weather affects supply and demand. We explicitly drop `precipProbability` due to strong multicollinearity with `precipIntensity` ($r = 0.84$), and exclude `cloudCover` and `short_summary` because our EDA showed minimal impact on pricing.

Interaction terms

We create interaction terms between distance and service-related variables, particularly `distance × service class`, because our EDA revealed distinct "banding" patterns suggesting different service tiers charge different rates per mile. We also include `distance × cab_type` to capture how platform choice and surge pricing affect marginal costs.

```
In [27]: numerical_cols = ['distance', 'surge_multiplier', 'temperature',
                     'precipIntensity', 'hour', 'is_weekend']

# combine everything
X_train_combined = pd.concat([
    X_train[numerical_cols].reset_index(drop=True),
    X_train_loc_encoded.reset_index(drop=True),
    X_train_categorical.reset_index(drop=True)
], axis=1)

X_val_combined = pd.concat([
    X_val[numerical_cols].reset_index(drop=True),
    X_val_loc_encoded.reset_index(drop=True),
    X_val_categorical.reset_index(drop=True)
], axis=1)

X_test_combined = pd.concat([
    X_test[numerical_cols].reset_index(drop=True),
    X_test_loc_encoded.reset_index(drop=True),
    X_test_categorical.reset_index(drop=True)
], axis=1)

print(f" Total features: {X_train_combined.shape[1]}")
print(f" - Numerical features: {len(numerical_cols)}")
print(f" - Target-encoded locations: {X_train_loc_encoded.shape[1]}")
print(f" - One-hot encoded categoricals: {X_train_categorical.shape[1]})")
```

```
Total features: 21
- Numerical features: 6
- Target-encoded locations: 2
- One-hot encoded categoricals: 13
```

Interaction terms

```
In [28]: X_train_final = X_train_combined.copy()
X_val_final = X_val_combined.copy()
X_test_final = X_test_combined.copy()

# distance x service name
interaction_count = 0
service_name_cols = [col for col in X_train_categorical.columns if 'name_' in col]

for col in service_name_cols:
    interaction_name = f'distance_x_{col}'
    X_train_final[interaction_name] = X_train_final['distance'] * X_train_final[col]
    X_val_final[interaction_name] = X_val_final['distance'] * X_val_final[col]
    X_test_final[interaction_name] = X_test_final['distance'] * X_test_final[col]
    interaction_count += 1

# distance x cab_type
cab_type_cols = [col for col in X_train_categorical.columns if 'cab_type_' in col]
for col in cab_type_cols:
    interaction_name = f'distance_x_{col}'
    X_train_final[interaction_name] = X_train_final['distance'] * X_train_final[col]
    X_val_final[interaction_name] = X_val_final['distance'] * X_val_final[col]
    X_test_final[interaction_name] = X_test_final['distance'] * X_test_final[col]
    interaction_count += 1

print(f"Created {interaction_count} interaction terms")
interaction_cols = [col for col in X_train_final.columns if '_x_' in col]
print("\nInteraction terms:")
for i, col in enumerate(interaction_cols, 1):
    print(f" {i}. {col}")
```

Created 13 interaction terms

Interaction terms:

1. distance_x_name_Black SUV
2. distance_x_name_Lux
3. distance_x_name_Lux Black
4. distance_x_name_Lux Black XL
5. distance_x_name_Lyft
6. distance_x_name_Lyft XL
7. distance_x_name_Shared
8. distance_x_name_UberPool
9. distance_x_name_UberX
10. distance_x_name_UberXL
11. distance_x_name_WAV
12. distance_x_month_name_November
13. distance_x_cab_type_Uber

Train multiple linear regression

```
In [29]: mlr_model = LinearRegression()
mlr_model.fit(X_train_final, y_train)

print(f"Features: {X_train_final.shape[1]}")
print(f"Intercept: {mlr_model.intercept_:.4f}")

y_train_pred_mlr = mlr_model.predict(X_train_final)
y_val_pred_mlr = mlr_model.predict(X_val_final)
y_test_pred_mlr = mlr_model.predict(X_test_final)

def calc_metrics(y_true, y_pred, set_name):
    r2 = r2_score(y_true, y_pred)
    rmse = np.sqrt(mean_squared_error(y_true, y_pred))
    print(f"{set_name}:")
    print(f" R²: {r2:.4f}")
    print(f" RMSE: {rmse:.4f}")
    return {'R²': r2, 'RMSE': rmse}

train_metrics = calc_metrics(y_train, y_train_pred_mlr, "Training")
val_metrics = calc_metrics(y_val, y_val_pred_mlr, "Validation")
```

Features: 34
 Intercept: 2.0154
 Training:
 R²: 0.9421
 RMSE: 0.1366
 Validation:
 R²: 0.9425
 RMSE: 0.1369

Multiple linear regression results

Our multiple linear regression model with 34 features achieves an R^2 of 0.9425 on the validation set, compared to just 0.1169 for the baseline distance-only model. This represents an 8-fold improvement in variance explained and a 74.5% reduction in RMSE (from 0.5368 to 0.1369). The near-identical performance on training ($R^2 = 0.9421$) and validation sets shows excellent generalization without overfitting. These results confirm that ride-sharing prices are determined not just by distance, but by a complex interplay of service class, platform, location, temporal patterns, and weather conditions.

```
In [30]: fig, axes = plt.subplots(1, 2, figsize=(16, 6))

# Training set
axes[0].scatter(y_train, y_train_pred_mlr, alpha=0.3, s=1, color='steelblue')
axes[0].plot([y_train.min(), y_train.max()],
            [y_train.min(), y_train.max()],
            'r--', lw=2, label='Perfect Prediction')
axes[0].set_xlabel('Actual log(Price)', fontsize=12)
axes[0].set_ylabel('Predicted log(Price)', fontsize=12)
axes[0].set_title(f'Training Set\nR2 = {train_metrics["R2"]:.4f}, RMSE = {train_metrics["RMSE"]:.4f}',
                 fontsize=14, fontweight='bold')
axes[0].legend()
axes[0].grid(True, alpha=0.3)

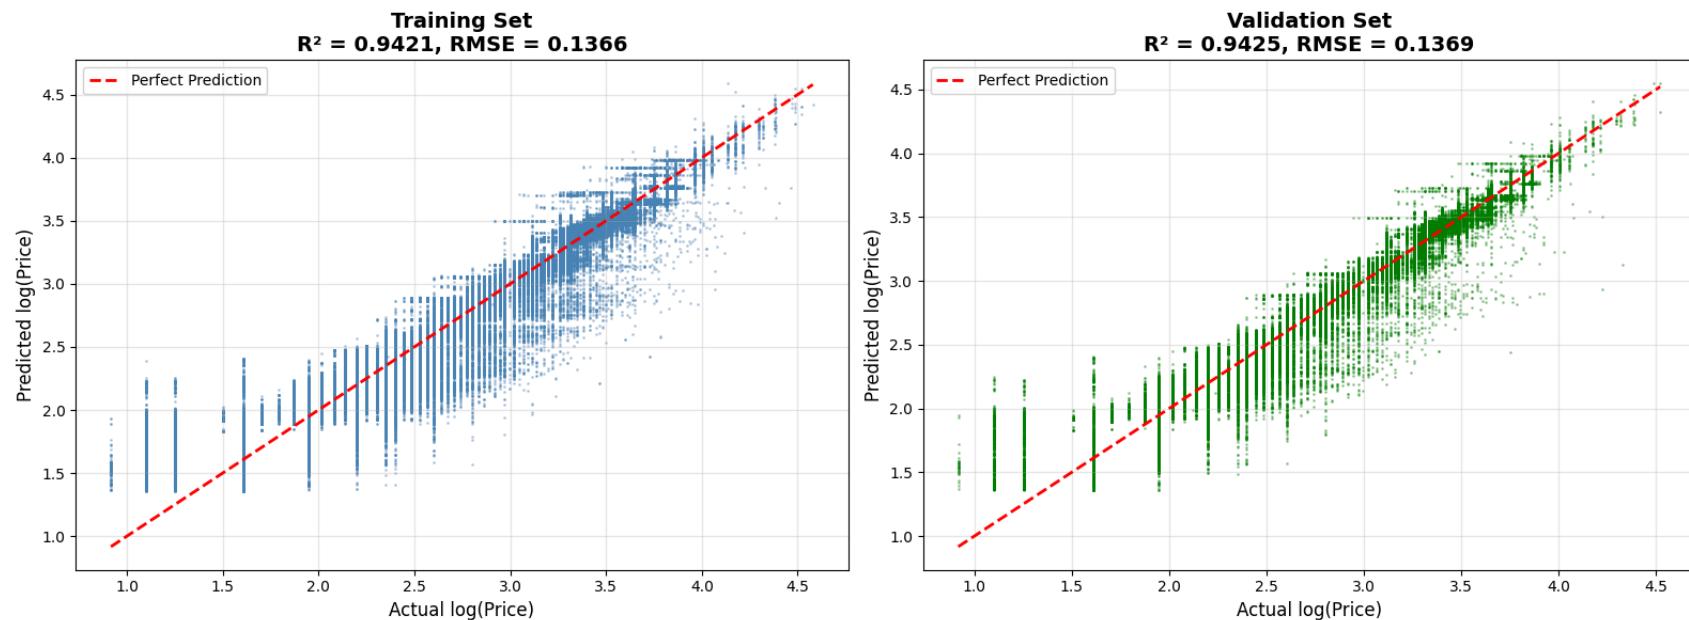
# Validation set
axes[1].scatter(y_val, y_val_pred_mlr, alpha=0.3, s=1, color='green')
axes[1].plot([y_val.min(), y_val.max()],
            [y_val.min(), y_val.max()],
            'r--', lw=2, label='Perfect Prediction')
axes[1].set_xlabel('Actual log(Price)', fontsize=12)
axes[1].set_ylabel('Predicted log(Price)', fontsize=12)
axes[1].set_title(f'Validation Set\nR2 = {val_metrics["R2"]:.4f}, RMSE = {val_metrics["RMSE"]:.4f}',
```

```

    fontsize=14, fontweight='bold')
axes[1].legend()
axes[1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

```



Feature importance

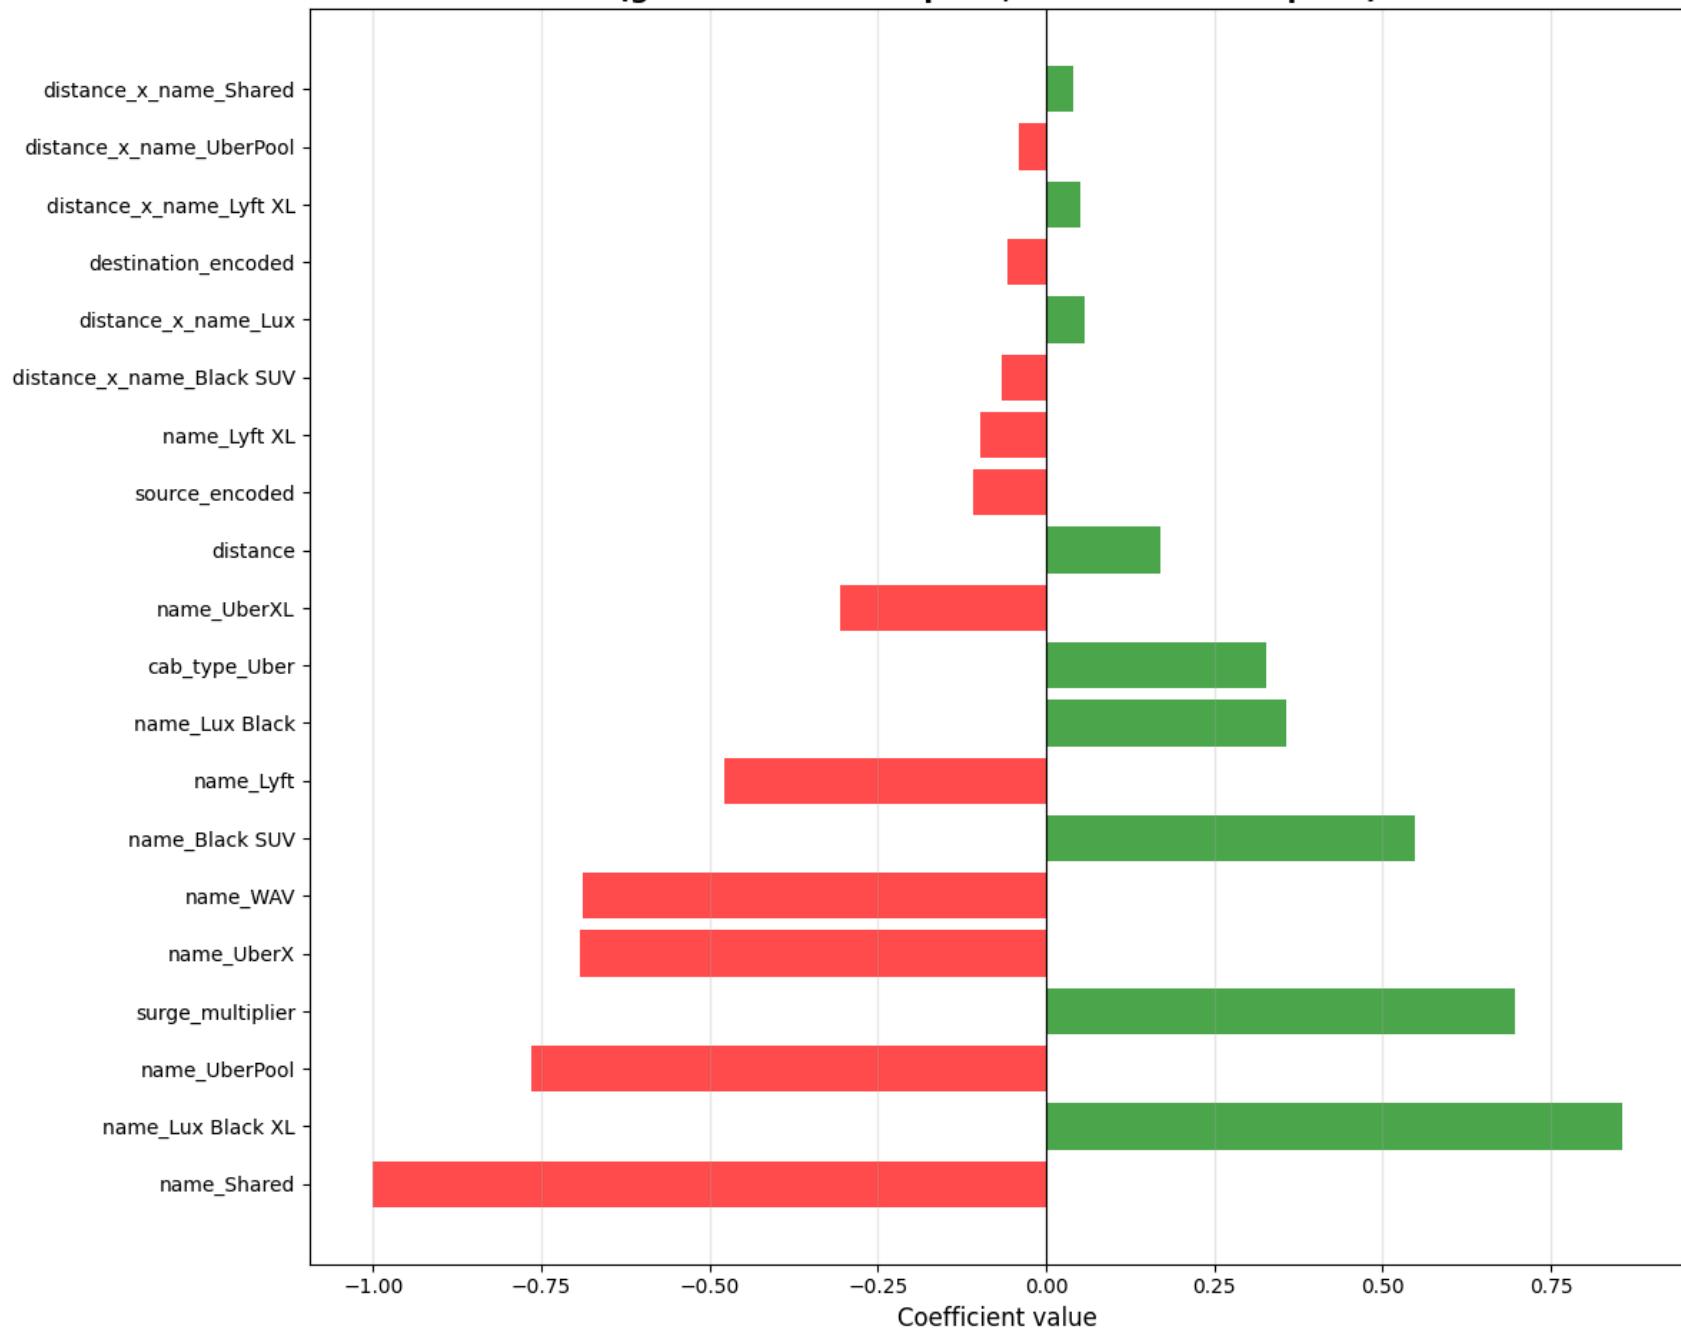
The coefficient analysis confirms that `service class` is the dominant pricing factor in the Boston ride-sharing market. Premium options like Lux Black XL (+0.85) and Black SUV (+0.52) command substantial premiums, while economy services like Shared rides (-1.05) and UberPool (-0.72) offer significant discounts. `Surge multiplier` (+0.70) emerges as the second most influential factor. `Distance` maintains a positive effect (+0.15), with interaction terms showing that premium services charge higher per-mile rates than economy options. The platform indicator shows `Uber` charges moderately more than `Lyft` (+0.33), while location effects (`source` and `destination` encodings) demonstrate neighborhood-level price variation.

```
In [31]: # top 20 features by absolute coefficient value
coef_df = pd.DataFrame({
```

```
'Feature': X_train_final.columns,
'Coefficient': mlr_model.coef_
})
coef_df['Abs_Coefficient'] = coef_df['Coefficient'].abs()
top_20_features = coef_df.nlargest(20, 'Abs_Coefficient')

fig, ax = plt.subplots(figsize=(12, 10))
colors = ['green' if x > 0 else 'red' for x in top_20_features['Coefficient']]
bars = ax.barh(range(len(top_20_features)), top_20_features['Coefficient'], color=colors, alpha=0.7)
ax.set_yticks(range(len(top_20_features)))
ax.set_yticklabels(top_20_features['Feature'], fontsize=10)
ax.set_xlabel('Coefficient value', fontsize=12)
ax.set_title('Top 20 most influential features\n(green = increases price, red = decreases price)', fontsize=14, fontweight='bold')
ax.axvline(x=0, color='black', linestyle='-', linewidth=0.8)
ax.grid(True, alpha=0.3, axis='x')
plt.tight_layout()
plt.show()
```

Top 20 most influential features
(green = increases price, red = decreases price)



Decision Tree Regression

While linear regression provides interpretable coefficients and strong performance, decision trees can capture non-linear relationships and complex interactions without explicitly defining them.

Thus, we'll build a decision tree model using cross-validation to optimize hyperparameters and prevent overfitting.

We use the **same data split** as MLR:

- Same target variable: `y_train` and `y_val`
- Same feature matrices: `X_train_final` and `X_val_final` (derived from the same initial split)
- Same random seed: `random_state=42` ensures reproducibility

```
In [32]: warnings.filterwarnings('ignore')

# define hyperparameter grid for cross-validation
param_grid = {
    'max_depth': [5, 10, 15, 20, 25, 30],
    'min_samples_split': [2, 5, 10, 20],
    'min_samples_leaf': [1, 2, 4, 8],
    'max_features': ['sqrt', 'log2', None]
}

# initialize decision tree regressor
dt_base = DecisionTreeRegressor(random_state=42)

# perform 5-fold cross-validation
print("Performing 5-fold cross-validation to find optimal hyperparameters...")
print("This may take a few minutes...\n")

grid_search = GridSearchCV(
    estimator=dt_base,
    param_grid=param_grid,
    cv=5,
    scoring='neg_root_mean_squared_error',
    n_jobs=-1,
    verbose=1
```

```
)  
  
# fit on training data  
grid_search.fit(X_train_final, y_train)  
  
print(f"\nBest hyperparameters found:")  
for param, value in grid_search.best_params_.items():  
    print(f" {param}: {value}")  
print(f"\nBest cross-validation RMSE: {-grid_search.best_score_:.4f}")
```

Performing 5-fold cross-validation to find optimal hyperparameters...
This may take a few minutes...

Fitting 5 folds for each of 288 candidates, totalling 1440 fits

Best hyperparameters found:

```
max_depth: 15  
max_features: None  
min_samples_leaf: 8  
min_samples_split: 20
```

Best cross-validation RMSE: 0.1302

```
In [33]: # train final model with best hyperparameters  
dt_model = grid_search.best_estimator_  
  
# make predictions  
y_train_pred_dt = dt_model.predict(X_train_final)  
y_val_pred_dt = dt_model.predict(X_val_final)  
y_test_pred_dt = dt_model.predict(X_test_final)  
  
# evaluate performance using same metrics as MLR  
train_metrics_dt = calc_metrics(y_train, y_train_pred_dt, "Training")  
val_metrics_dt = calc_metrics(y_val, y_val_pred_dt, "Validation")  
test_metrics_dt = calc_metrics(y_test, y_test_pred_dt, "Test")
```

```
Training:  
R2: 0.9553  
RMSE: 0.1201  
Validation:  
R2: 0.9485  
RMSE: 0.1296  
Test:  
R2: 0.9486  
RMSE: 0.1289
```

Decision tree results

The optimized decision tree model shows strong predictive performance while maintaining good generalization. Cross-validation identified the optimal tree configuration to balance model complexity with prediction accuracy, preventing both underfitting and overfitting.

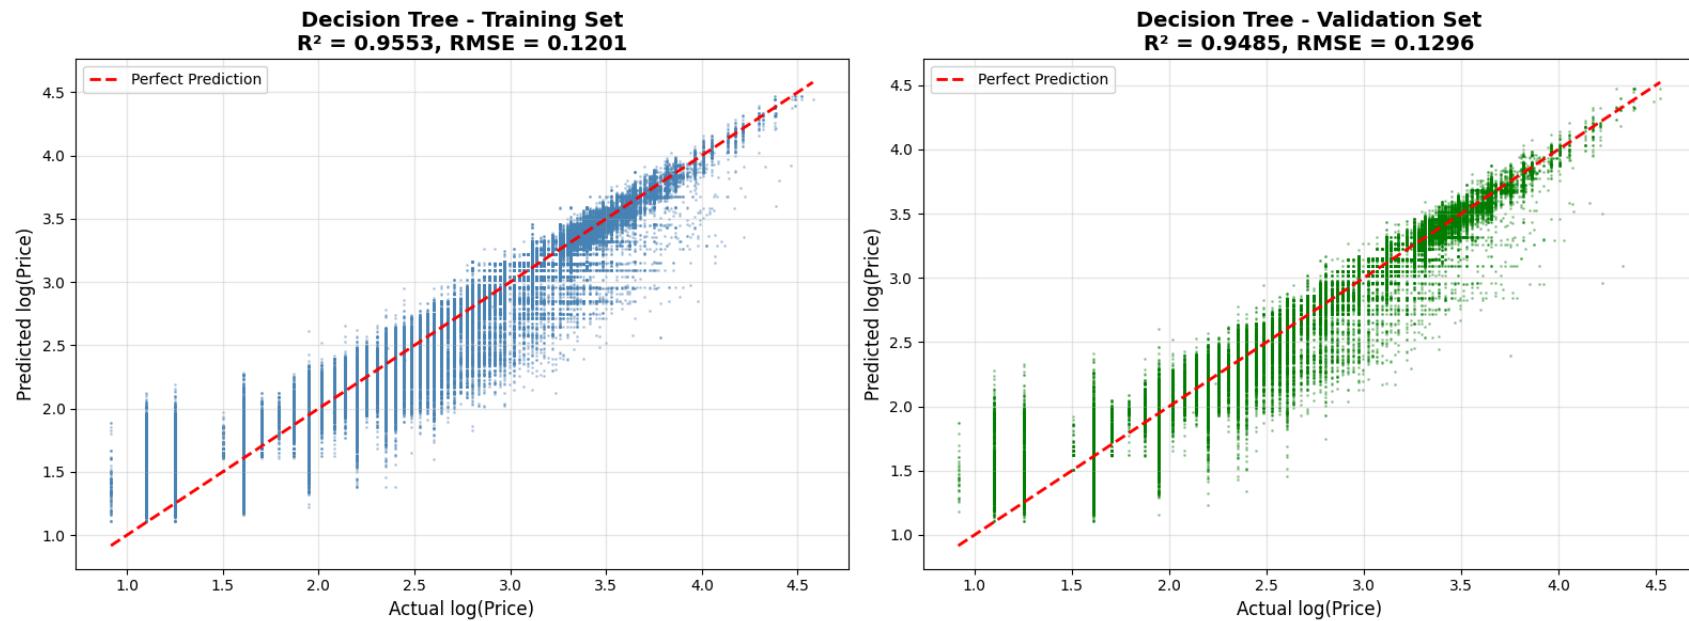
```
In [34]: # visualization: predicted vs actual (same style as MLR)  
fig, axes = plt.subplots(1, 2, figsize=(16, 6))  
  
# Training set  
axes[0].scatter(y_train, y_train_pred_dt, alpha=0.3, s=1, color='steelblue')  
axes[0].plot([y_train.min(), y_train.max()],  
            [y_train.min(), y_train.max()],  
            'r--', lw=2, label='Perfect Prediction')  
axes[0].set_xlabel('Actual log(Price)', fontsize=12)  
axes[0].set_ylabel('Predicted log(Price)', fontsize=12)  
axes[0].set_title(f'Decision Tree - Training Set\nR2 = {train_metrics_dt["R2"]:.4f}, RMSE = {train_metrics_dt["RMSE"]:.4f}',  
                 fontsize=14, fontweight='bold')  
axes[0].legend()  
axes[0].grid(True, alpha=0.3)  
  
# Validation set  
axes[1].scatter(y_val, y_val_pred_dt, alpha=0.3, s=1, color='green')  
axes[1].plot([y_val.min(), y_val.max()],  
            [y_val.min(), y_val.max()],  
            'r--', lw=2, label='Perfect Prediction')  
axes[1].set_xlabel('Actual log(Price)', fontsize=12)  
axes[1].set_ylabel('Predicted log(Price)', fontsize=12)  
axes[1].set_title(f'Decision Tree - Validation Set\nR2 = {val_metrics_dt["R2"]:.4f}, RMSE = {val_metrics_dt["RMSE"]:.4f}',  
                 fontsize=14, fontweight='bold')
```

```

axes[1].legend()
axes[1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

```



Feature importance

Decision trees provide feature importance measures based on the **mean decrease in Gini impurity**. Each feature's importance is calculated as the total reduction in node impurity (weighted by the probability of reaching that node) when that feature is used for splitting across all nodes in the tree. This directly reflects how much each feature contributes to reducing variance in the predictions.

```

In [35]: # extract feature importances
feature_importance_df = pd.DataFrame({
    'Feature': X_train_final.columns,
    'Importance': dt_model.feature_importances_
})

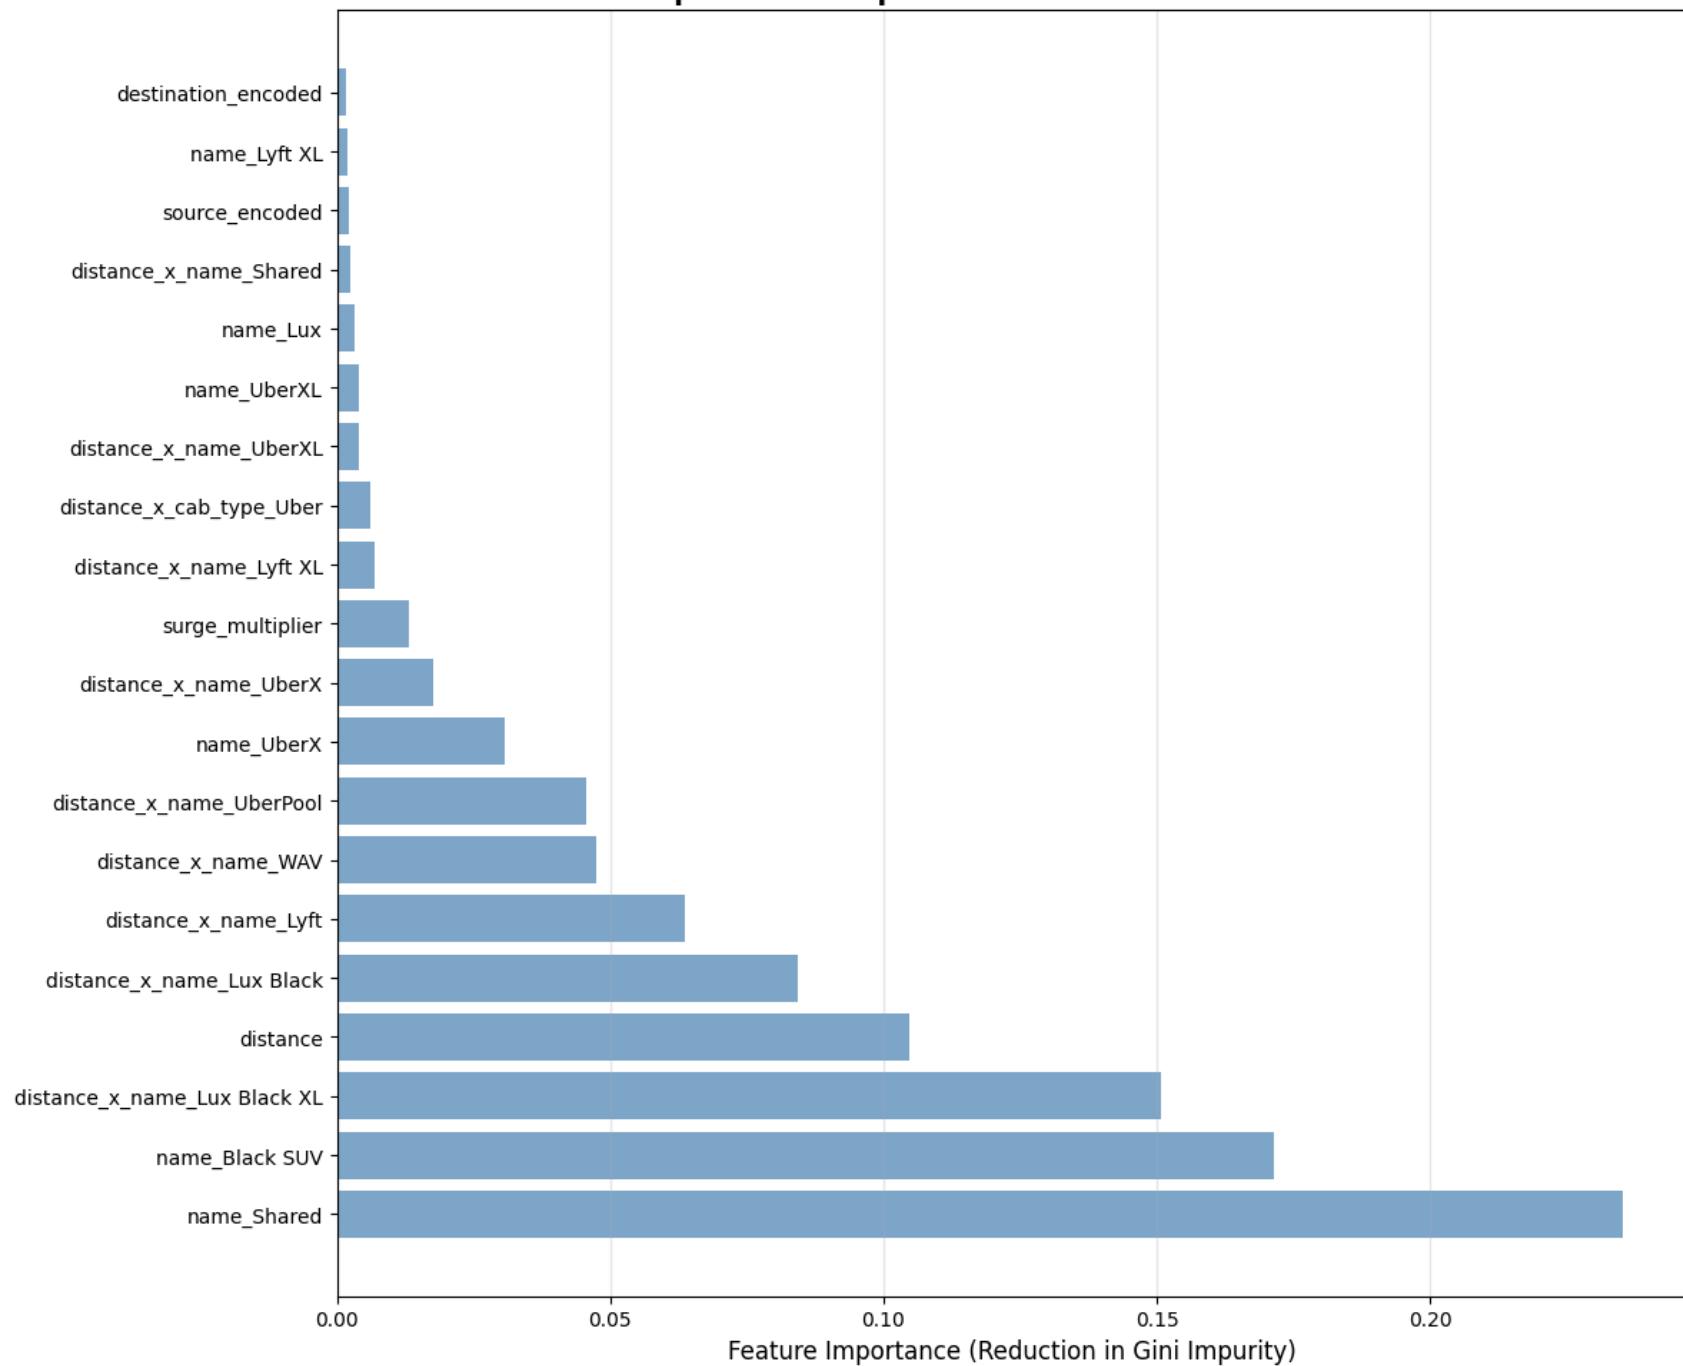
# get top 20 features

```

```
top_20_dt_features = feature_importance_df.nlargest(20, 'Importance')

# plot feature importances
fig, ax = plt.subplots(figsize=(12, 10))
colors_dt = ['steelblue' for _ in range(len(top_20_dt_features))]
bars = ax.barh(range(len(top_20_dt_features)), top_20_dt_features['Importance'], color=colors_dt, alpha=0.7)
ax.set_yticks(range(len(top_20_dt_features)))
ax.set_yticklabels(top_20_dt_features['Feature'], fontsize=10)
ax.set_xlabel('Feature Importance (Reduction in Gini Impurity)', fontsize=12)
ax.set_title('Top 20 Most Important Features - Decision Tree',
             fontsize=14, fontweight='bold')
ax.grid(True, alpha=0.3, axis='x')
plt.tight_layout()
plt.show()

print("\nTop 20 Feature Importances:")
for i, (idx, row) in enumerate(top_20_dt_features.iterrows(), 1):
    print(f"{i:2d}. {row['Feature']}:40s {row['Importance']:.4f}")
```

Top 20 Most Important Features - Decision Tree

Top 20 Feature Importances:

1. name_Shared	0.2352
2. name_Black SUV	0.1715
3. distance_x_name_Lux Black XL	0.1507
4. distance	0.1048
5. distance_x_name_Lux Black	0.0843
6. distance_x_name_Lyft	0.0636
7. distance_x_name_WAV	0.0475
8. distance_x_name_UberPool	0.0456
9. name_UberX	0.0306
10. distance_x_name_UberX	0.0176
11. surge_multiplier	0.0132
12. distance_x_name_Lyft XL	0.0068
13. distance_x_cab_type_Uber	0.0062
14. distance_x_name_UberXL	0.0041
15. name_UberXL	0.0041
16. name_Lux	0.0032
17. distance_x_name_Shared	0.0025
18. source_encoded	0.0022
19. name_Lyft XL	0.0021
20. destination_encoded	0.0016

Model comparison: Decision Tree vs Multiple Linear Regression

Comparing the decision tree with our multiple linear regression model reveals the trade-offs between model flexibility and interpretability.

```
In [36]: # create comparison table
comparison_df = pd.DataFrame({
    'Model': ['Multiple Linear Regression', 'Decision Tree'],
    'Train R2': [train_metrics['R2'], train_metrics_dt['R2']],
    'Train RMSE': [train_metrics['RMSE'], train_metrics_dt['RMSE']],
    'Validation R2': [val_metrics['R2'], val_metrics_dt['R2']],
    'Validation RMSE': [val_metrics['RMSE'], val_metrics_dt['RMSE']],
    'Test R2': [np.nan, test_metrics_dt['R2']],
    'Test RMSE': [np.nan, test_metrics_dt['RMSE']]
})

print("Model Performance Comparison:")
print("=*90)
```

```
print(comparison_df.to_string(index=False))
print("=*90)

# calculate overfitting gap
mlr_gap = train_metrics['R2'] - val_metrics['R2']
dt_gap = train_metrics_dt['R2'] - val_metrics_dt['R2']

print(f"\nOverfitting Analysis (Train R2 - Validation R2):")
print(f"  MLR: {mlr_gap:.4f}")
print(f"  Decision Tree: {dt_gap:.4f}")

# visualize comparison
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# R2 comparison
models = ['MLR', 'Decision Tree']
train_r2 = [train_metrics['R2'], train_metrics_dt['R2']]
val_r2 = [val_metrics['R2'], val_metrics_dt['R2']]

x = np.arange(len(models))
width = 0.35

axes[0].bar(x - width/2, train_r2, width, label='Training', color='steelblue', alpha=0.8)
axes[0].bar(x + width/2, val_r2, width, label='Validation', color='green', alpha=0.8)
axes[0].set_ylabel('R2 Score', fontsize=12)
axes[0].set_title('R2 Score Comparison', fontsize=14, fontweight='bold')
axes[0].set_xticks(x)
axes[0].set_xticklabels(models)
axes[0].legend()
axes[0].grid(True, alpha=0.3, axis='y')
axes[0].set_ylim([0.85, 1.0])

# RMSE comparison
train_rmse = [train_metrics['RMSE'], train_metrics_dt['RMSE']]
val_rmse = [val_metrics['RMSE'], val_metrics_dt['RMSE']]

axes[1].bar(x - width/2, train_rmse, width, label='Training', color='steelblue', alpha=0.8)
axes[1].bar(x + width/2, val_rmse, width, label='Validation', color='green', alpha=0.8)
axes[1].set_ylabel('RMSE', fontsize=12)
axes[1].set_title('RMSE Comparison', fontsize=14, fontweight='bold')
axes[1].set_xticks(x)
axes[1].set_xticklabels(models)
```

```

axes[1].legend()
axes[1].grid(True, alpha=0.3, axis='y')

plt.tight_layout()
plt.show()

```

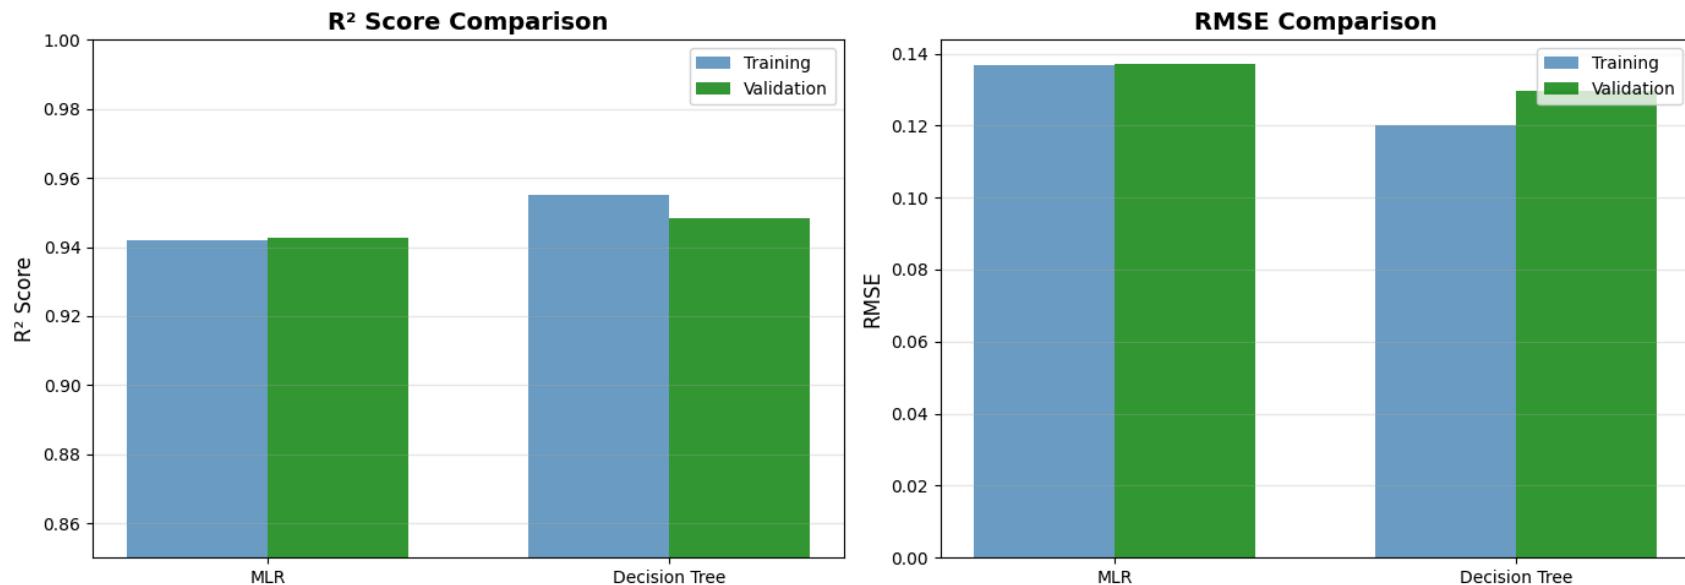
Model Performance Comparison:

	Model	Train R ²	Train RMSE	Validation R ²	Validation RMSE	Test R ²	Test RMSE
Multiple Linear Regression	MLR	0.942102	0.136648	0.942526	0.136949	NaN	NaN
	Decision Tree	0.955254	0.120130	0.948510	0.129624	0.94863	0.128937

Overfitting Analysis (Train R² - Validation R²):

MLR: -0.0004

Decision Tree: 0.0067



Summary: MLR vs DECISION TREE

The decision tree model slightly outperforms multiple linear regression, achieving a validation R² of 0.9485 compared to MLR's 0.9425, with lower RMSE (0.1296 vs. 0.1369). Both models show excellent generalization with minimal overfitting—the

decision tree has a train-validation R² gap of only 0.0067, while MLR actually generalizes slightly better to validation data (gap of -0.0004).

The key advantage of the decision tree is its ability to automatically capture non-linear relationships and complex interactions without manual feature engineering. The feature importance analysis reveals that service class indicators (name_Shared, name_Black SUV) and distance-based interactions dominate predictions, consistent with MLR's coefficient analysis. However, the decision tree identifies these patterns through recursive partitioning rather than linear coefficients.

While both models achieve strong predictive performance ($R^2 > 0.94$), MLR offers superior interpretability through its coefficients, making it easier to quantify the marginal effect of each feature. The decision tree provides more flexibility but at the cost of interpretability. For this pricing prediction task, either model would be suitable, with the choice depending on whether interpretability (favor MLR) or capturing complex non-linearities (favor decision tree) is prioritized.

Polynomial Regression

Build Polynomial Regression (Distance-Only Model)

Since the distance-price relationship is non-linear, we want to improve over the distance-only baseline by adding polynomial terms. Extending the baseline model, we will start with degree 2-3 (most common), evaluate each and compare.

```
In [37]: # Distance-only features for polynomial regression
X_train_dist = X_train[['distance']].copy()
X_val_dist = X_val[['distance']].copy()
X_test_dist = X_test[['distance']].copy()

def evaluate_poly_model(degree):
    # create polynomial transformer
    poly = PolynomialFeatures(degree=degree, include_bias=False)

    # transform distance-only input
    X_train_poly = poly.fit_transform(X_train_dist)
    X_val_poly = poly.transform(X_val_dist)

    # fit model
```

```
model = LinearRegression()
model.fit(X_train_poly, y_train)

# predictions
train_pred = model.predict(X_train_poly)
val_pred = model.predict(X_val_poly)

# metrics
r2_train = r2_score(y_train, train_pred)
r2_val = r2_score(y_val, val_pred)

rmse_train = np.sqrt(mean_squared_error(y_train, train_pred))
rmse_val = np.sqrt(mean_squared_error(y_val, val_pred))

print(f"\nPolynomial Degree {degree}")
print(f" Train R2: {r2_train:.4f}, RMSE: {rmse_train:.4f}")
print(f" Val R2: {r2_val:.4f}, RMSE: {rmse_val:.4f}")

return {
    'degree': degree,
    'train_r2': r2_train,
    'val_r2': r2_val,
    'train_rmse': rmse_train,
    'val_rmse': rmse_val
}

results_poly = []

for d in [1, 2, 3]:
    results_poly.append(evaluate_poly_model(d))
```

Polynomial Degree 1
Train R²: 0.1135, RMSE: 0.5347
Val R²: 0.1169, RMSE: 0.5368

Polynomial Degree 2
Train R²: 0.1160, RMSE: 0.5340
Val R²: 0.1194, RMSE: 0.5361

Polynomial Degree 3
Train R²: 0.1161, RMSE: 0.5339
Val R²: 0.1196, RMSE: 0.5360

```
In [38]: dist_plot = np.linspace(0, X_train['distance'].max(), 500).reshape(-1, 1)

plt.figure(figsize=(12, 6))

colors = {
    1: 'red',      # degree 1 (baseline linear)
    2: 'forestgreen',  # degree 2
    3: 'gold'        # degree 3
}

for degree in [1, 2, 3]:
    poly = PolynomialFeatures(degree=degree, include_bias=False)
    dist_poly = poly.fit_transform(dist_plot)

    model = LinearRegression()
    model.fit(poly.fit_transform(X_train_dist), y_train)

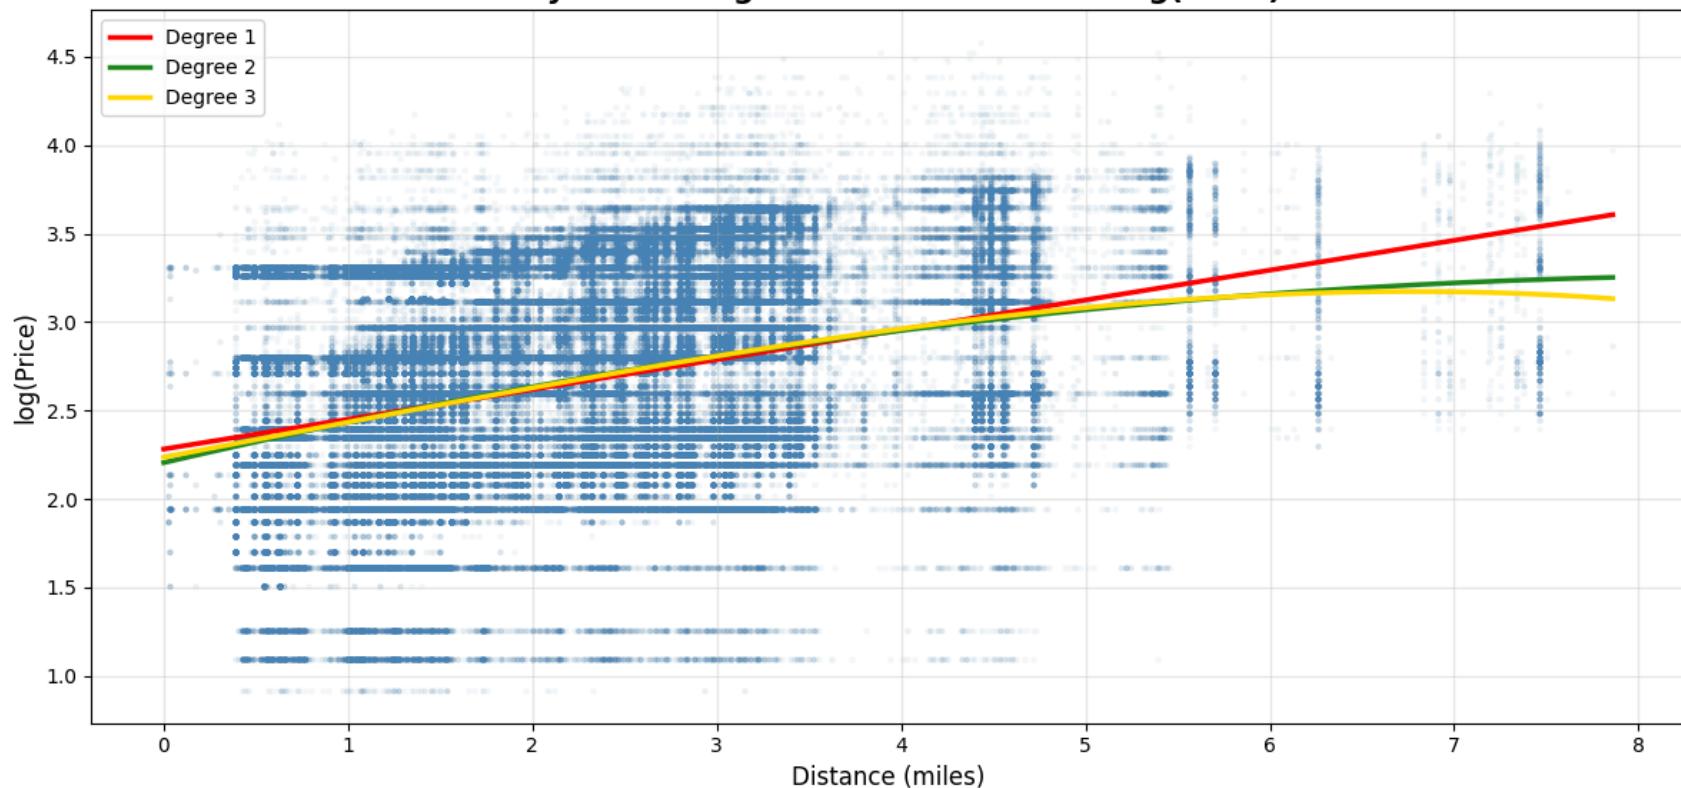
    plt.plot(
        dist_plot,
        model.predict(dist_poly),
        label=f'Degree {degree}',
        linewidth=2.5,
        color=colors[degree]
    )

plt.scatter(
    X_train['distance'], y_train,
    alpha=0.03, s=5,
    color='steelblue'
)

plt.xlabel("Distance (miles)", fontsize=12)
plt.ylabel("log(Price)", fontsize=12)
plt.title("Polynomial Regression – Distance vs log(Price)", fontsize=15, fontweight='bold')

plt.legend(frameon=True)
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```

Polynomial Regression - Distance vs log(Price)



Polynomial regression improves the distance-only model slightly, but the effect is marginal. The distance-price relationship is nearly linear, and higher-degree polynomial terms cannot compensate for the absence of key categorical and contextual features. This motivates the need for more complex models incorporating service type, surge, location, and time, which is exactly what we explore in the full MLR and decision tree models.

Polynomial Regression with Selected Features

The MLR results showed that service class, surge multiplier, platform (Uber vs. Lyft), and distance, particularly through its interactions with service tiers, are the primary drivers of $\log(\text{price})$, while weather and temporal variables play secondary roles. Building on these findings, the polynomial regression incorporates this core subset of important features while introducing nonlinear transformations of distance (distance^2 and distance^3) to test whether allowing curvature in the distance–price relationship improves predictive performance beyond the linear specification.

We will use the features that previous MLR section identified as most important:

- distance
- surge_multiplier
- temperature
- precipIntensity
- hour
- is_weekend
- cab_type
- name (service class)
- distance²
- distance³

```
In [39]: # Selected numerical features
numerical_cols = [
    'distance', 'surge_multiplier', 'temperature',
    'precipIntensity', 'hour', 'is_weekend'
]

# Selected categorical features
categorical_cols = ['cab_type', 'name']

# Extract base data
X_train_s2 = X_train[numerical_cols + categorical_cols].copy()
X_val_s2 = X_val[numerical_cols + categorical_cols].copy()
X_test_s2 = X_test[numerical_cols + categorical_cols].copy()

# Dummy encoding
X_train_s2 = pd.get_dummies(X_train_s2, columns=categorical_cols, drop_first=True)
X_val_s2 = pd.get_dummies(X_val_s2, columns=categorical_cols, drop_first=True)
X_test_s2 = pd.get_dummies(X_test_s2, columns=categorical_cols, drop_first=True)

# Ensure same columns across splits
X_val_s2 = X_val_s2.reindex(columns=X_train_s2.columns, fill_value=0)
X_test_s2 = X_test_s2.reindex(columns=X_train_s2.columns, fill_value=0)

for df in [X_train_s2, X_val_s2, X_test_s2]:
    df['distance_sq'] = df['distance'] ** 2
```

```
df['distance_cu'] = df['distance'] ** 3

poly_s2_model = LinearRegression()
poly_s2_model.fit(X_train_s2, y_train)

# Predictions
y_train_pred_s2 = poly_s2_model.predict(X_train_s2)
y_val_pred_s2 = poly_s2_model.predict(X_val_s2)
y_test_pred_s2 = poly_s2_model.predict(X_test_s2)

def calc_metrics(y_true, y_pred, set_name=''):
    r2 = r2_score(y_true, y_pred)
    rmse = np.sqrt(mean_squared_error(y_true, y_pred))
    print(f"{set_name}:")
    print(f" R²: {r2:.4f}")
    print(f" RMSE: {rmse:.4f}")
    return {'R²': r2, 'RMSE': rmse}

# Evaluate
train_s2_metrics = calc_metrics(y_train, y_train_pred_s2, "Training (Poly S2)")
val_s2_metrics = calc_metrics(y_val, y_val_pred_s2, "Validation (Poly S2)")
test_s2_metrics = calc_metrics(y_test, y_test_pred_s2, "Test (Poly S2)")

fig, axes = plt.subplots(1, 2, figsize=(16, 6))

# Training set
axes[0].scatter(y_train, y_train_pred_s2, alpha=0.3, s=1, color='steelblue')
axes[0].plot([y_train.min(), y_train.max()],
            [y_train.min(), y_train.max()],
            'r--', lw=2, label='Perfect Prediction')
axes[0].set_xlabel('Actual log(Price)', fontsize=12)
axes[0].set_ylabel('Predicted log(Price)', fontsize=12)
axes[0].set_title(
    f'Polynomial Regression (Selected Features) – Training\n'
    f'R² = {train_s2_metrics["R²"]:.4f}, RMSE = {train_s2_metrics["RMSE"]:.4f}\n',
    fontsize=14, fontweight='bold'
)
axes[0].legend()
axes[0].grid(True, alpha=0.3)

# Validation set
axes[1].scatter(y_val, y_val_pred_s2, alpha=0.3, s=1, color='green')
```

```
axes[1].plot([y_val.min(), y_val.max()],
             [y_val.min(), y_val.max()],
             'r--', lw=2, label='Perfect Prediction')
axes[1].set_xlabel('Actual log(Price)', fontsize=12)
axes[1].set_ylabel('Predicted log(Price)', fontsize=12)
axes[1].set_title(
    f'Polynomial Regression (Selected Features) - Validation\n'
    f'R² = {val_s2_metrics["R²"]:.4f}, RMSE = {val_s2_metrics["RMSE"]:.4f}',
    fontsize=14, fontweight='bold'
)
axes[1].legend()
axes[1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```

Training (Poly S2):

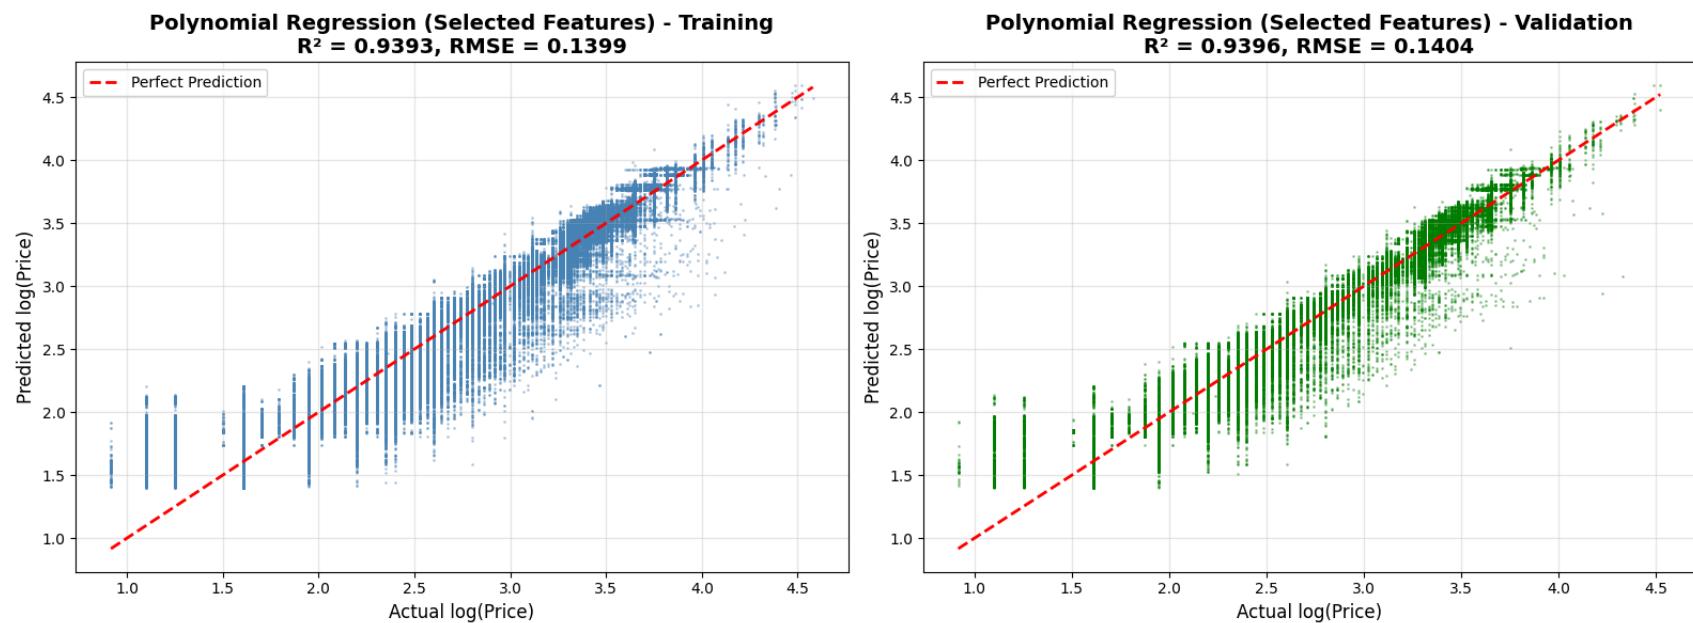
R²: 0.9393
RMSE: 0.1399

Validation (Poly S2):

R²: 0.9396
RMSE: 0.1404

Test (Poly S2):

R²: 0.9398
RMSE: 0.1396



These results represent a substantial improvement over the baseline distance-only models (R^2 is about 0.12) and show that incorporating service class, surge, platform, and temporal features dramatically increases explanatory power. Compared to the full MLR ($R^2 = 0.9425$) and the decision tree ($R^2 = 0.9485$), the polynomial model performs slightly worse, indicating that allowing distance to be nonlinear adds only limited predictive value once the key categorical and contextual variables are present.